

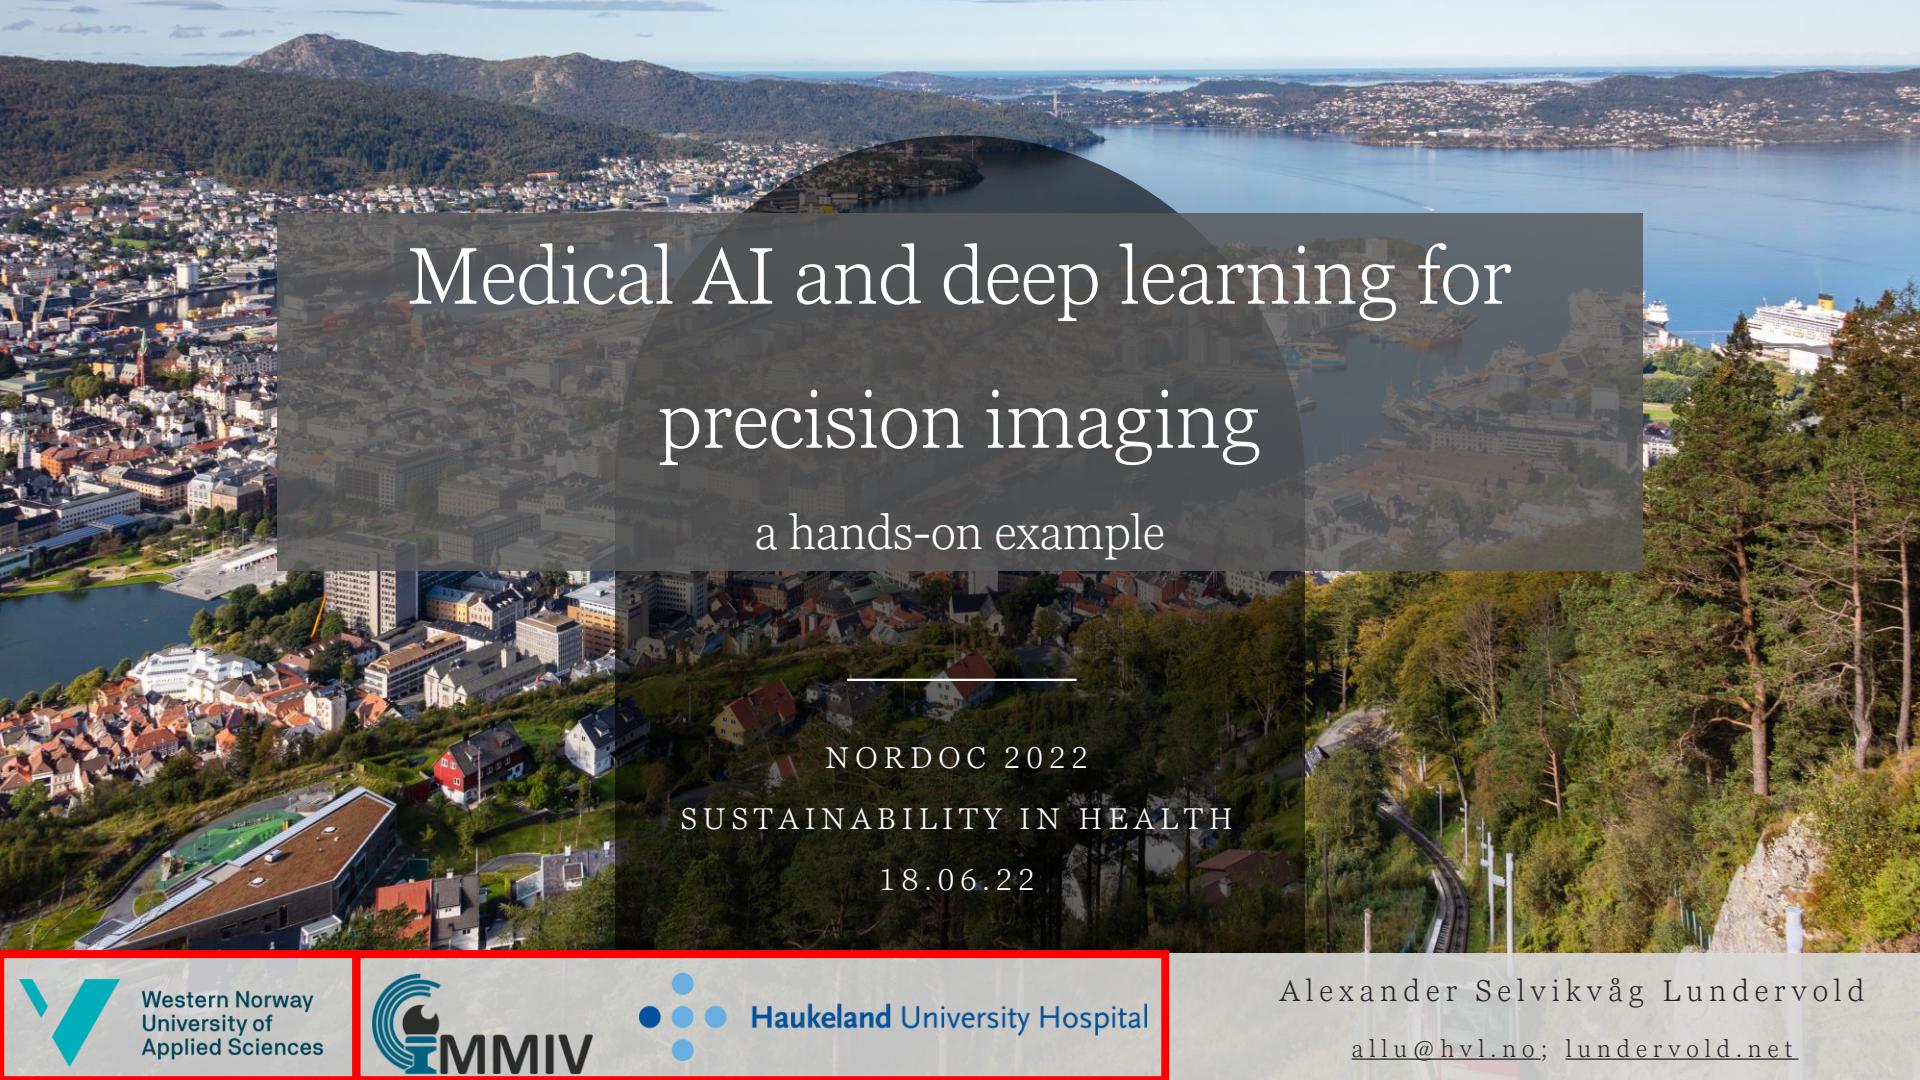
# Medical AI and deep learning for precision imaging

a hands-on example

NORDOC 2022  
SUSTAINABILITY IN HEALTH

18.06.22





# Medical AI and deep learning for precision imaging

a hands-on example

NORDOC 2022  
SUSTAINABILITY IN HEALTH

18.06.22

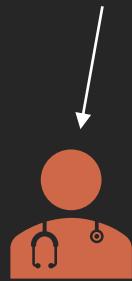
# Plan



Hands-on!



What *is* deep  
learning?



How do you *do*  
deep learning?



Some perspectives,  
opportunities and  
challenges

# Plan



What is deep learning?

Objectives:

- (i) to *demystify*
- (ii) make sure everyone *understands* what deep learning is about at a fundamental level, and how you *do* deep learning as a practitioner
- (iii) hopefully trigger some ideas and potentially a *desire to explore further*

How do you *do* deep learning?

Some perspectives,  
opportunities and  
challenges



Hands-on!

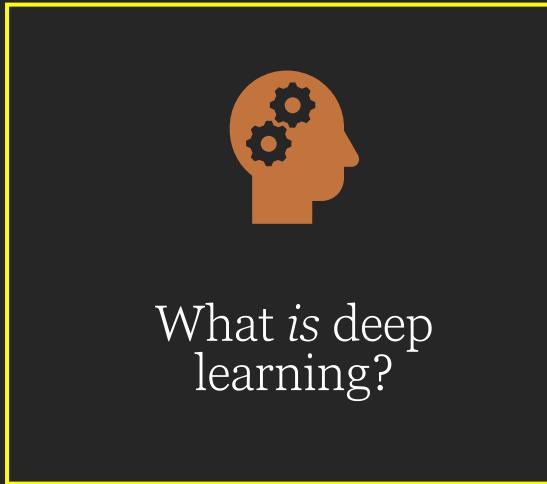
/

Will only have time for a quick peek at the field!

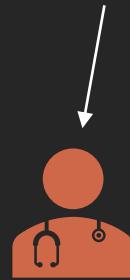
# Plan



Hands-on!



What *is* deep learning?



How do you *do* deep learning?



Some perspectives,  
opportunities and  
challenges

**Deep learning**

*searching for good hierarchical geometric representations*

## Deep learning

*searching for good hierarchical geometric representations*

## Deep learning

*searching for **good** hierarchical geometric representations*

## Deep learning

*searching for good **hierarchical** geometric representations*

**Deep learning**

*searching for good hierarchical geometric representations*

01

02

03

04

05

Function approximation

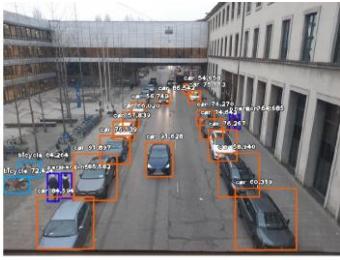
Function approximation

$$y \approx f(x; \theta)$$

what's in the image

an image

The diagram illustrates the concept of function approximation in machine learning. It features a mathematical equation  $y \approx f(x; \theta)$ . Two arrows point from descriptive text to specific parts of the equation: one arrow points from the phrase "what's in the image" to the input variable  $x$ , and another arrow points from the phrase "an image" to the output variable  $y$ .

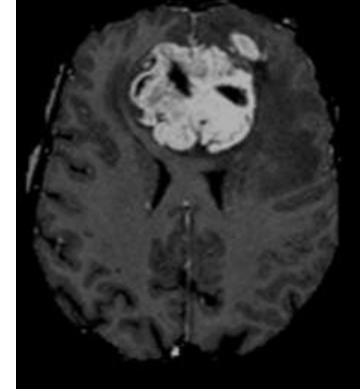
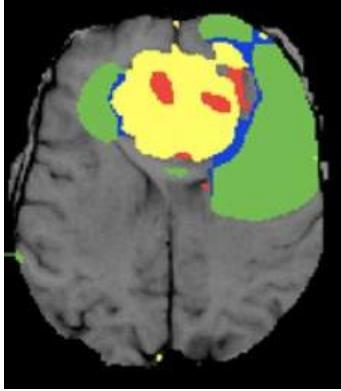


Function approximation

$$y \approx f(x; \theta)$$

what's in the image

an image



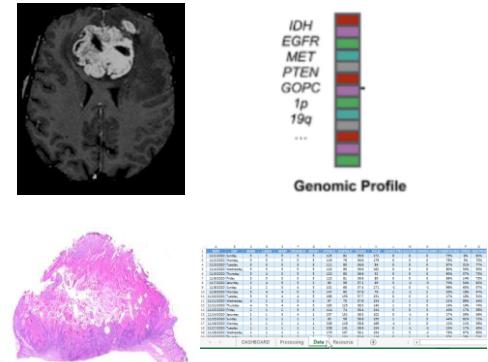
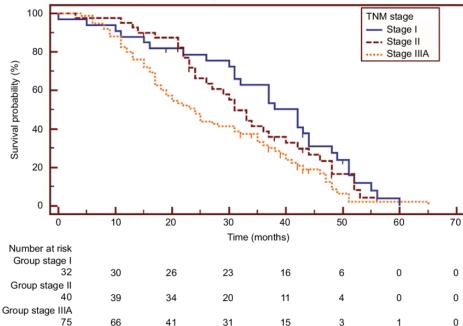
Function approximation

$$y \approx f(x; \theta)$$

what's in the image

an image

Arrows point from the text "what's in the image" to the yellow cluster in the brain segmentation map, and from the text "an image" to the grayscale brain scan.



## Function approximation

$$y \approx f(x; \theta)$$

outcome (e.g., survival)

heterogeneous information

Training

$$(X, y)$$

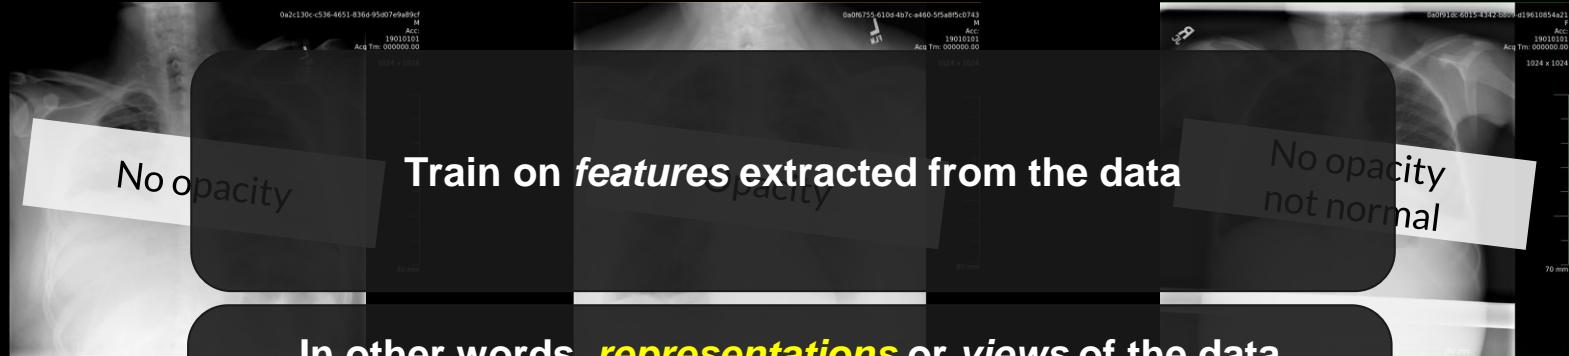
## Training

$$(X, y)$$

Input data

Labels

*This is the setup in what's called **supervised learning***



In other words, **representations** or **views** of the data

*Important that these representations are sufficiently informative*



$$y \approx f(x ; \theta)$$

# 01

## Machine learning: **function approximation** based on **training**

$$y \approx f(x ; \theta)$$

healthy or not healthy /  
outcome

features of a disease,  
process, patient or  
system

## Models and parameters

$$y \approx f(x; \theta)$$


The diagram consists of two black arrows originating from the words "model" and "parameters" at the bottom. The arrow from "model" points to the left side of the function  $f$ . The arrow from "parameters" points to the right side of the function  $f$ .

***Training*** is a search for useful representations through a space defined by a family  $\mathcal{F}$  of parametrized models

**01**

Machine learning: **function approximation** based on **training**

**02**

We pick the model family  $\mathcal{F}$ . The parameters  $\Theta$  are found automatically through training

$$y \approx f(x; \theta)$$

## Theory versus practice

Theory       $X \longrightarrow Y$

**In theory:**

can get arbitrarily good approximations by training simple models  
(*universal approximation* and *no free lunch theorems*)

Practice      $X \xrightarrow{\text{complex function}} Y$

**In practice:**

which  $f$  we use and how the training is done makes a huge difference

The family  $F$  of models considered, i.e., the *hypothesis space*, should be

- (i) **efficiently searchable**,
- (ii) leading to **expressive** models that can
- (iii) **generalize** beyond the training data distribution (i.e., work well on new data)

**01**

Machine learning: **function approximation** based on **training**

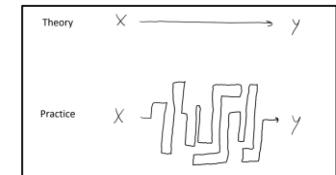
**02**

We pick the model family  $\mathcal{F}$ . The parameters  $\Theta$  are found automatically through training

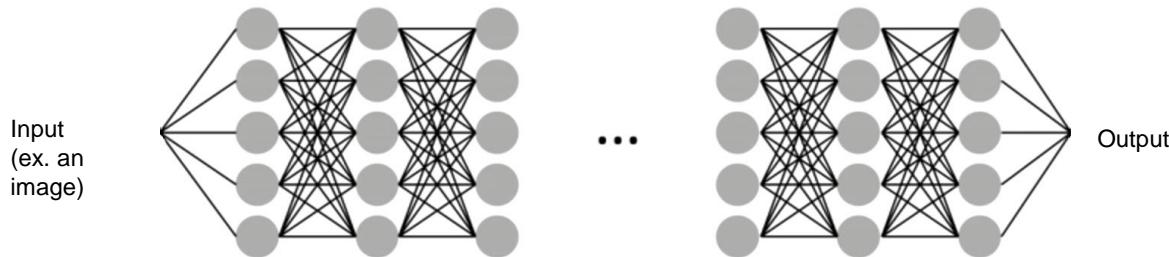
**03**

In practice: some models are better suited to some tasks than others.

$$y \approx f(x; \theta)$$

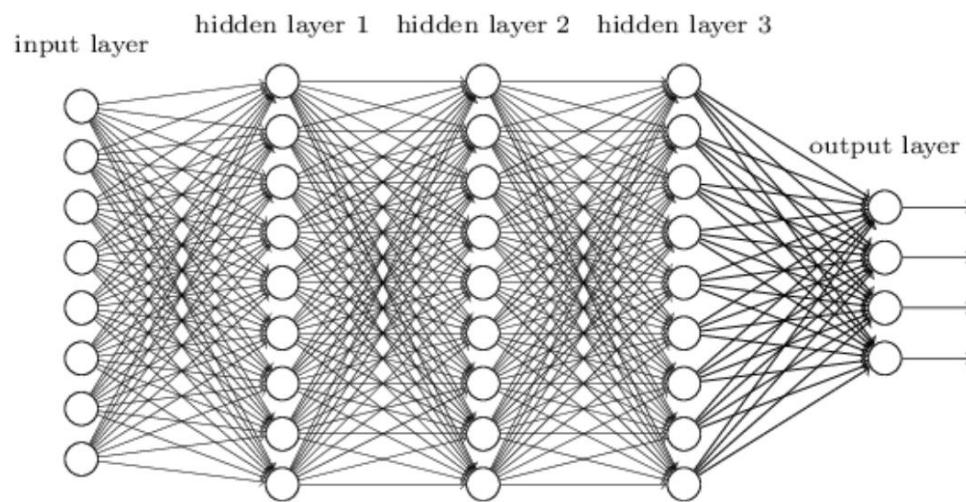


$$y \approx f(x; \theta)$$

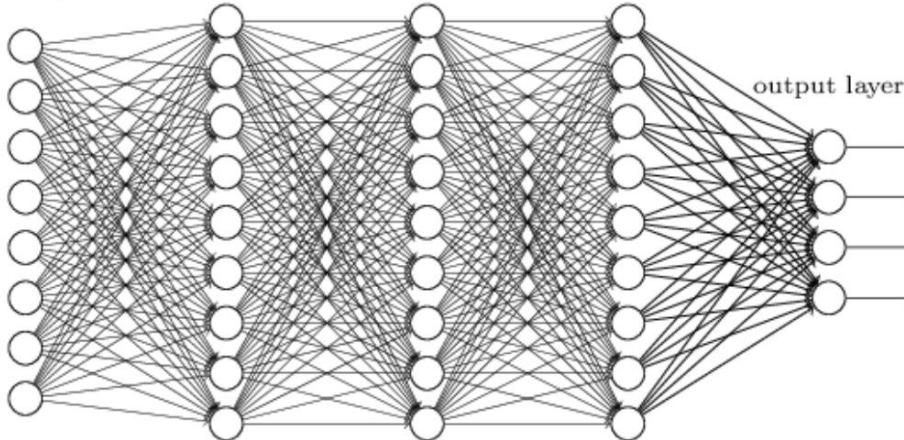


### Artificial neural networks

Computational graphs of simple units (“neurons”) connected in various specific ways, parametrized by parameters associated to each layer.



input layer      hidden layer 1    hidden layer 2    hidden layer 3



A geometric analogy:  
uncrumpling paper



This is a fancy way to draw what's really just a composition of simple functions:

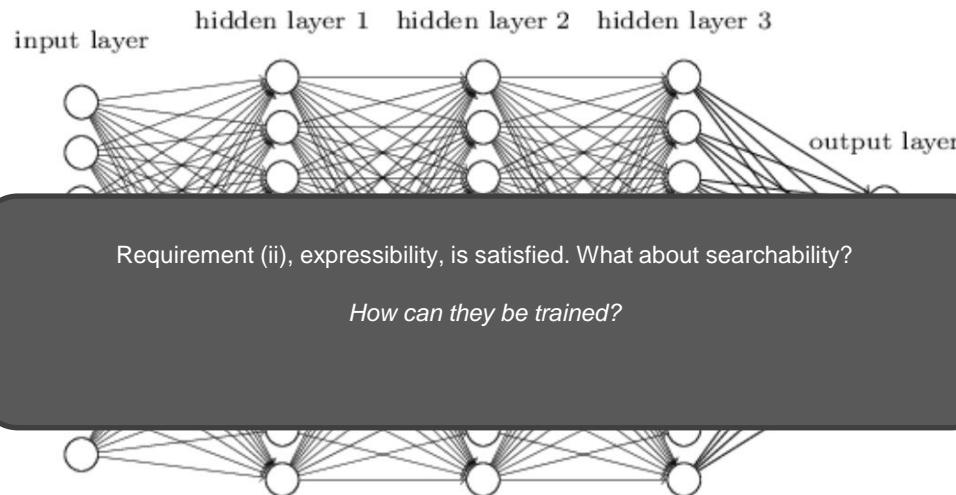
$$y = \sigma_{k+1}(W_{k+1}^T \sigma_k(W_k^T \left( \dots \left( \sigma_1(W_1^T x) \right) \dots \right)) = f(x; \theta)$$

Input  $x$  is sent through several consecutive **layers** (functions). An output is produced.

Think of each layer as a (geometric) **transformation** of its input.

Every layer has a bunch of **parameters**, stored in the matrices  $W$

*Each family member of  $F$  is complex transformation broken down into a composition of simpler ones,  
parametrized by the layer weights.*



This is a fancy way to draw what's really just a composition of simple functions:

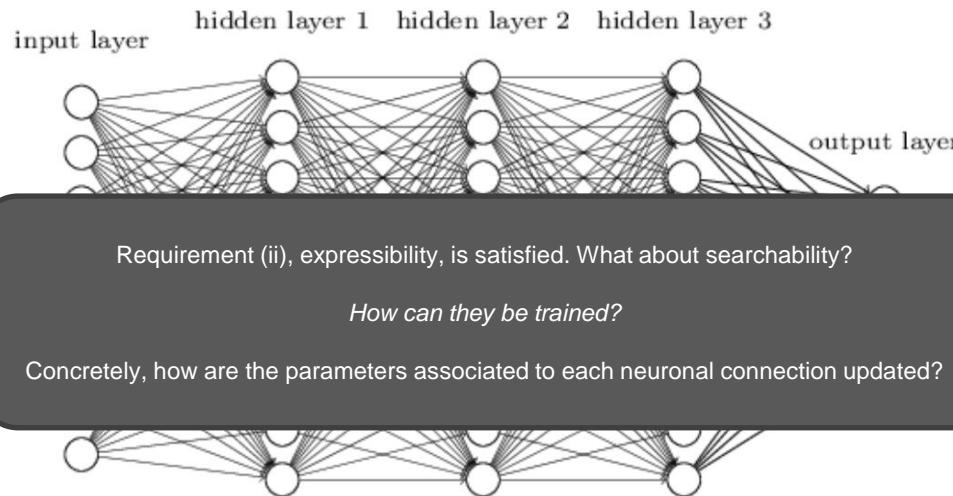
$$y = \sigma_{k+1}(W_{k+1}^T \sigma_k(W_k^T \left( \dots \left( \sigma_1(W_1^T x) \right) \dots \right)) = f(x; \theta)$$

Input  $x$  is sent through several consecutive **layers** (functions). An output is produced.

Think of each layer as a (geometric) **transformation** of its input.

Every layer has a bunch of **parameters**, stored in the matrices  $W$

*Each family member of  $F$  is complex transformation broken down into a composition of simpler ones, parametrized by the layer weights.*



This is a fancy way to draw what's really just a composition of simple functions:

$$y = \sigma_{k+1}(W_{k+1}^T \sigma_k(W_k^T \left( \dots \left( \sigma_1(W_1^T x) \right) \dots \right)) = f(x; \theta)$$

Input  $x$  is sent through several consecutive **layers** (functions). An output is produced.

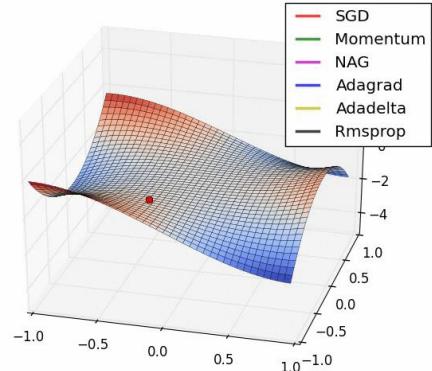
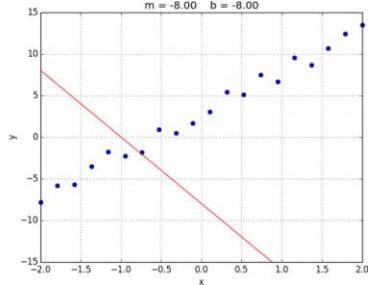
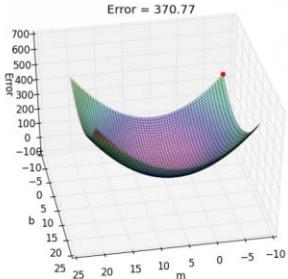
Think of each layer as a (geometric) **transformation** of its input.

Every layer has a bunch of **parameters**, stored in the matrices  $W$

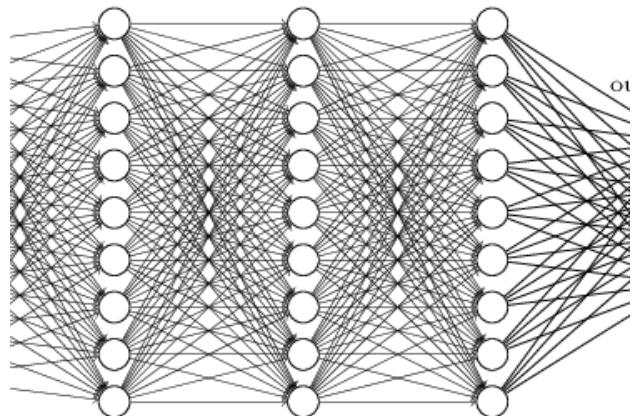
*Each family member of  $F$  is complex transformation broken down into a composition of simpler ones, parametrized by the layer weights.*

# Three basic ingredients:

1. A **loss function**: Used to measure the network's performance on the training data. A **feedback signal** for the training process.
2. An **optimizer**: Used to update the *parameters* of the network to increase performance as measured by the loss function.  
**Gradient descent and backpropagation.**
3. One or more **metrics**: A way to score the model. For example *accuracy* or *mean squared error*.



hidden layer 1   hidden layer 2   hidden layer 3

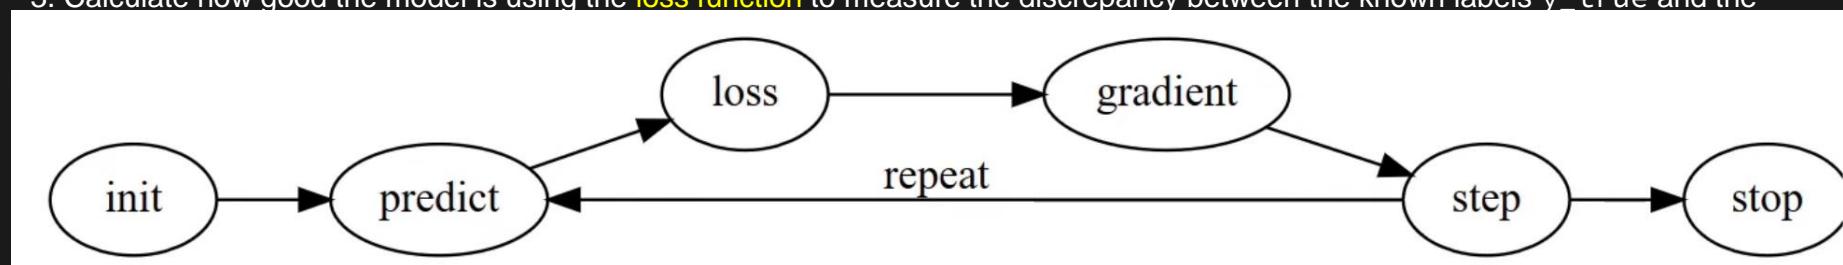


Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize the parameters
2. Collect a batch of inputs ( $X, y$ ) and use the model with its current parameters to make predictions  $y_{pred}$
3. Calculate how good the model is using the loss function to measure the discrepancy between the known labels  $y_{true}$  and the predictions  $y_{pred}$  for this batch
4. Calculate the gradients for each parameter using backpropagation. This measures how changing each parameter would change the loss. In other words, each parameter's contribution to the loss.
5. Move each weight a little bit in the opposite direction of its gradient using gradient descent. I.e., some variant of  
$$\text{parameter} = \text{parameter} - \text{learning\_rate} * \text{gradient}$$
6. Go back to step 2 and repeat the process
7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria

Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize the parameters
2. Collect a batch of inputs ( $X, y$ ) and use the model with its current parameters to make predictions  $y_{pred}$
3. Calculate how good the model is using the loss function to measure the discrepancy between the known labels  $y_{true}$  and the predictions  $y_{pred}$
4. Move each weight a little bit in the opposite direction of its gradient using gradient descent. I.e., some variant of  
$$\text{parameter} = \text{parameter} - \text{learning\_rate} * \text{gradient}$$
5. Go back to step 2 and repeat the process
6. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria



Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize the parameters

2. Collect a batch of inputs ( $X, y$ ) and use the model with its current parameters to make predictions  $y_{pred}$

3. Calculate how good the model is using the loss function to measure the discrepancy between the known labels  $y_{true}$  and the predictions  $y_{pred}$  for this batch

## What is gradient descent?

4. Calculate the gradients for each parameter. A gradient is a vector pointing from the current point on the loss surface to where changing each parameter would change the loss. In other words, each parameter's contribution to the loss.

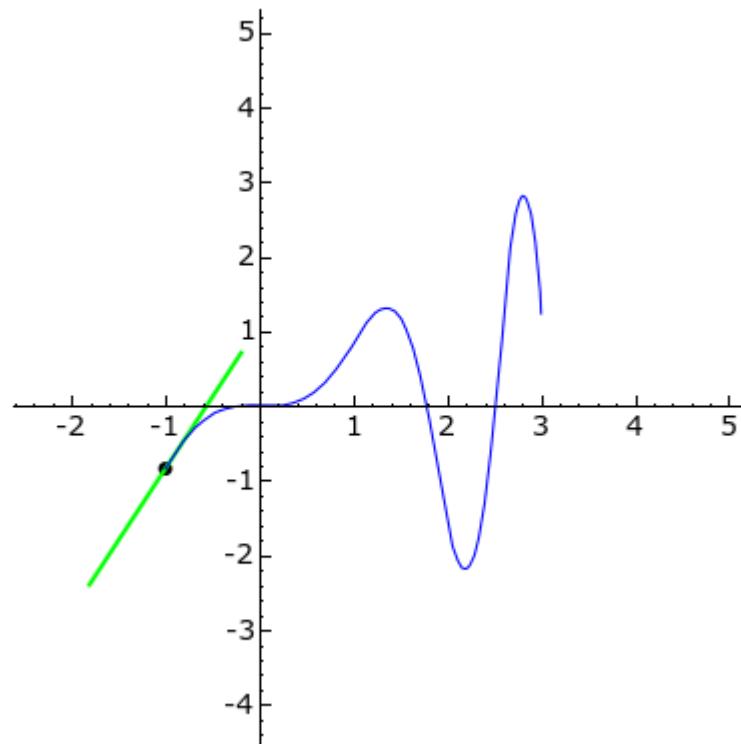
5. Move each weight a little bit in the opposite direction of its gradient using gradient descent. I.e., some variant of  
$$\text{parameter} = \text{parameter} - \text{learning\_rate} * \text{gradient}$$

6. Go back to step 2 and repeat the process

7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria

Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Init



2. Co

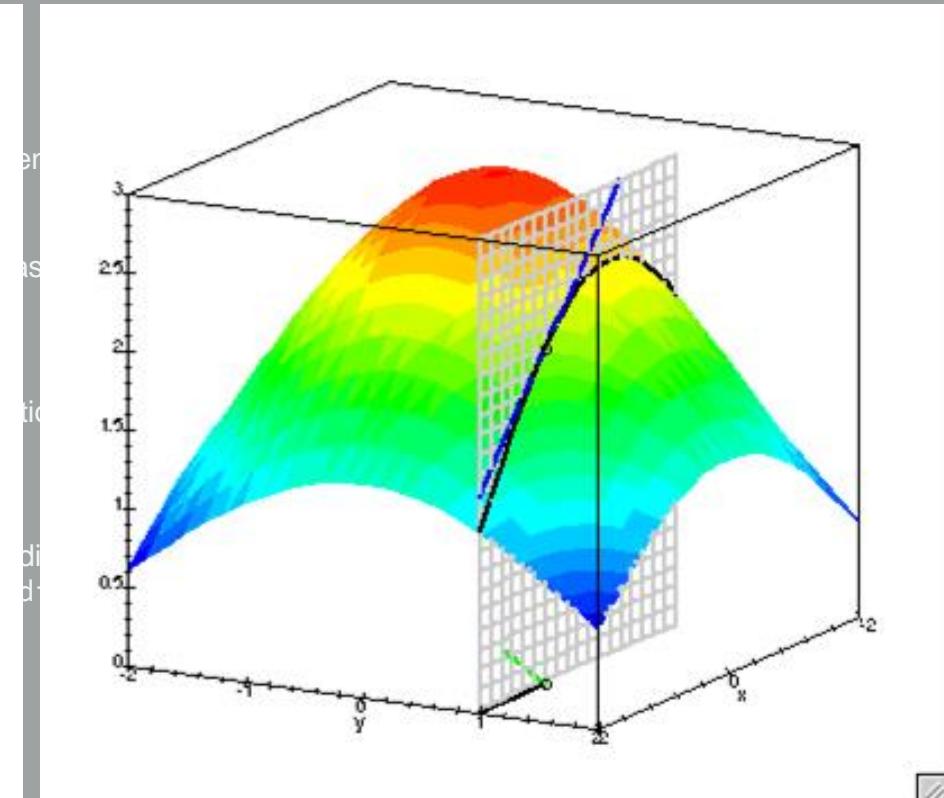
3. Ca

4. Ca

5. M

6. Go

7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria



Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Init

2. Co

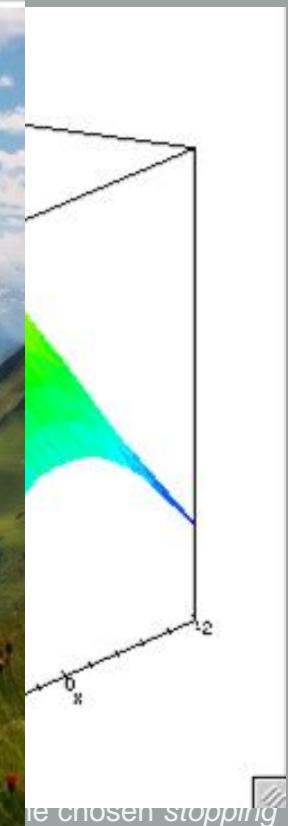
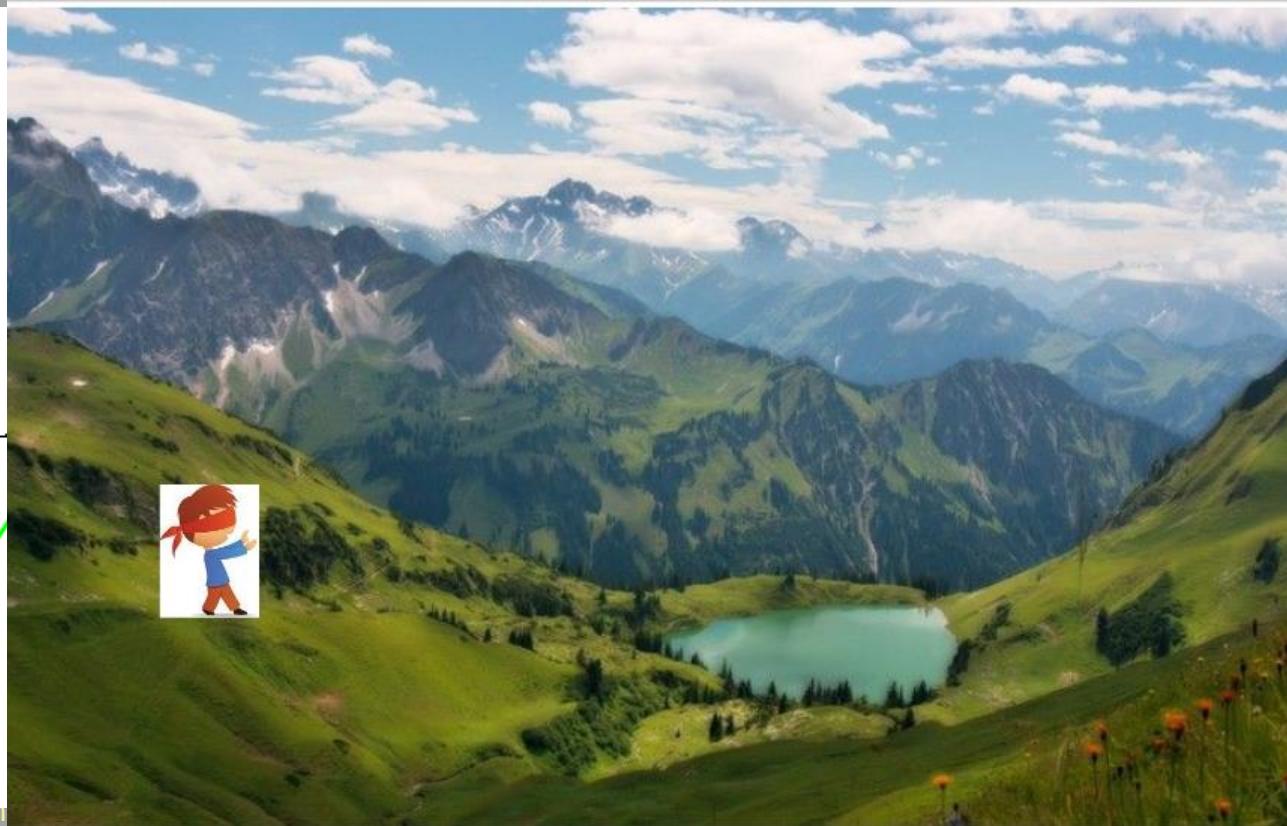
3. Ca  
pred

4. Ca  
loss.

5. M

6. Go

7. Do this over an  
criteria



the chosen stopping

**01**

Machine learning: **function approximation** based on **training**

**02**

We pick the model family  $F$ . The parameters  $\Theta$  are found automatically through training

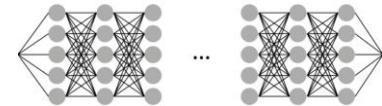
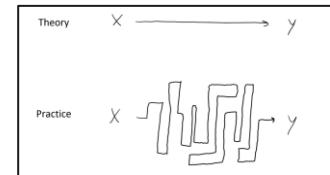
**03**

In practice: some models are better suited to some tasks than others.

**04**

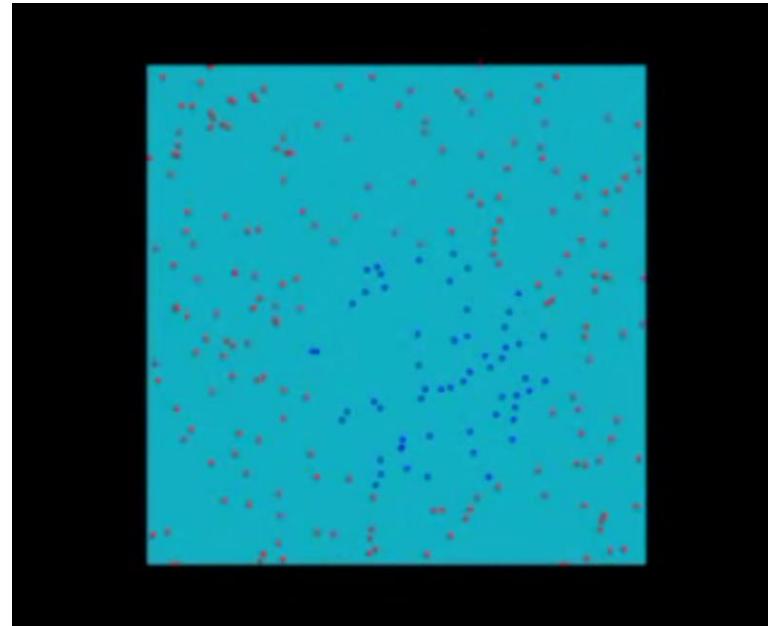
In deep learning: a particular choice of  $F$ . A **composition** of simple functions, trained jointly, typically using GPUs or other accelerators.

$$y \approx f(x; \theta)$$

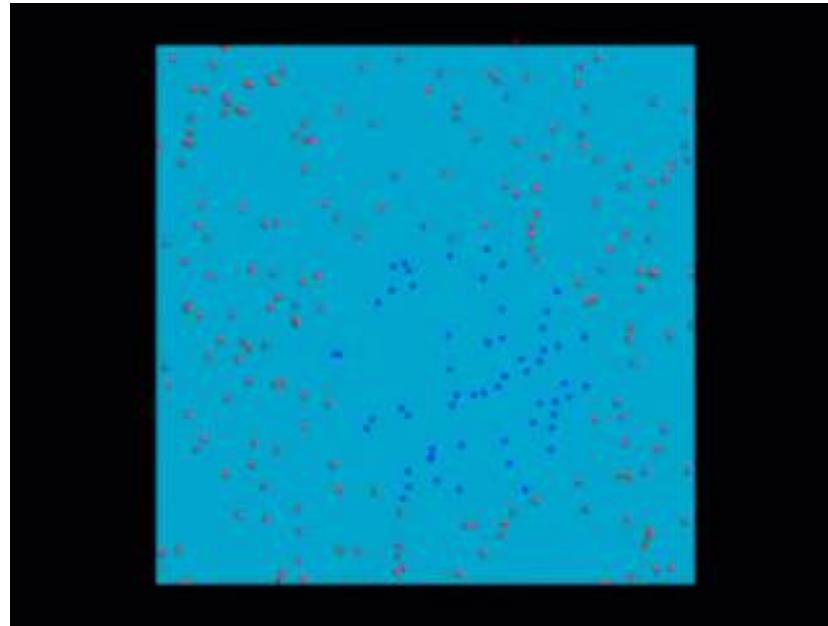


Representations and the power of transformations

## Transformations

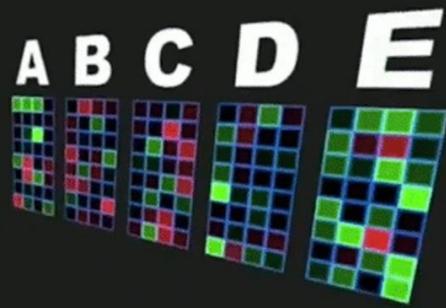


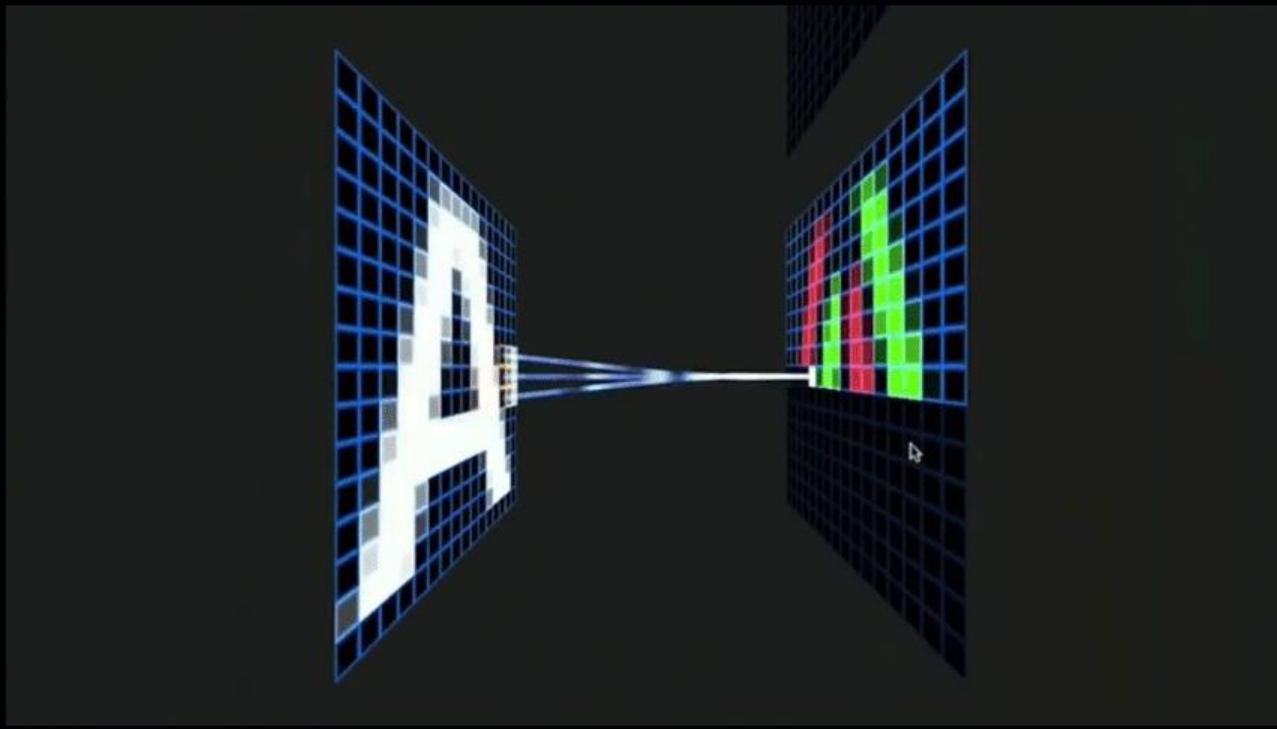
## Transformation



What transformations should be applied?

In deep learning: ***transformations providing useful representations automatically found via training***





**01**

Machine learning: **function approximation** based on **training**

**02**

We pick the model family  $F$ . The parameters  $\Theta$  are found automatically through training

**03**

In practice: some models are better suited to some tasks than others.

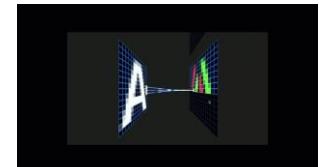
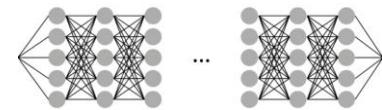
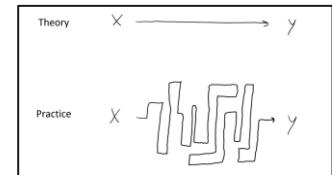
**04**

In deep learning: a particular choice of  $F$ . A **composition** of simple functions, trained jointly, typically using GPUs or other accelerators.

**05**

In deep learning: finding a **hierarchical set of geometric transformations**, ending in a representation that makes the desired task easy to solve.

$$y \approx f(x; \theta)$$



**01**

Machine learning: **function approximation** based on **training**

**02**

We pick the model family  $F$ . The parameters  $\Theta$  are found automatically through training

**03**

In practice: some models are better suited to some tasks than others.

### Deep learning

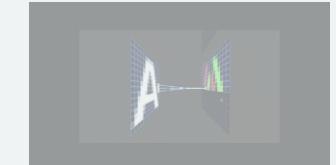
**04**

*searching for good hierarchical geometric representations*  
In deep learning: a particular choice of  $F$ . A composition of simple functions, trained jointly, typically using GPUs or other accelerators.

**05**

In deep learning: finding a **hierarchical set of geometric transformations**, ending in a representation that makes the desired task easy to solve.

$$y \approx f(x; \theta)$$



# Plan



What *is* deep  
learning?



How do you *do*  
deep learning?



Some perspectives,  
opportunities and  
challenges

jupyter 1.0-asl-segmentation-brain\_tumor\_segmentation Last Checkpoint: an hour ago (autosaved)

No kernel Not Trusted monai

ASL, 17.06.22

## Introduction

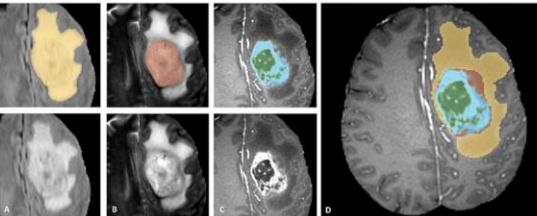
This notebook is based on the `brats_segmentation_3d.ipynb` notebook from <https://github.com/Project-MONAI/tutorials>. You should consult the MONAI documentation for a more thorough introduction to the MONAI framework: <https://monai.io>.

This notebook will give a hands-on example of deep learning in practice. Along the way, we will see some of the deep learning concepts covered in the introduction to the workshop translated into code.

Our goal is to illustrate a possible approach to brain tumor segmentation, i.e., extracting meaningful tumor sub-regions. This can lead to important imaging biomarkers such as f.ex. volumes and locations, and also provide regions of interest for extracting radiomics features. The next notebook will take a practical look at radiomics.

## The data: The Brain Tumor Segmentation (BraTS) challenge

We will use the data from the BraTS 2016 and 2017 competitions, consisting of 750 multi-parametric MRI studies of patients with brain tumors. Each study consists of four MRI modalities, T1w, contrast-enhanced T1w, T2w and FLAIR, and corresponding manual delineations of three tumor regions: necrotic tumor, active tumor and edema.



**Sample images**

jupyter 2.0-asl-brain\_tumor\_analysis\_radiomics Last Checkpoint: a minute ago (unsaved changes)

Trusted hdgio

## Introduction

In this notebook, we'll use a deep learning model to segment brain tumors from multi-parametric MRI and then extract features from the resulting tumor masks. Such features can potentially be associated with tumor severity and prognosis and contribute to better treatment.

We'll use data from the TCGA collection discussed by Arvid earlier today (see the slides at <https://github.com/MMLV-MI/nordic2022>):



**The Cancer Genome Atlas Program Data: X**

**TCGA-06-1802** data set including metadata  
 (DICOM images converted to NIFTI format using [dicom2nii](#))  
 TCGA-06-1802\_5\_AX\_T2\_FR-FSL.nii.gz  
 TCGA-06-1802\_5\_AX\_T2\_POST\_FSL.nii.gz  
 TCGA-06-1802\_5\_AX\_T2\_POST\_FSL.nii.gz  
 TCGA-06-1802\_A\_AX\_T2\_FLAIR.nii.gz  
 TCGA-06-1802\_7\_AX\_T1\_pre\_GD\_FLAIR.json  
 TCGA-06-1802\_7\_AX\_T1\_POST\_GD\_FLAIR.nii.gz  
 TCGA-06-1802\_8\_AX\_T1\_POST\_GD\_FLAIR.json  
 TCGA-06-1802\_R\_AX\_T1\_POST\_GD\_FLAIR.nii.gz  
 TCGA-06-1802\_clinical.tsv  
 TCGA-06-1802\_exposure.tsv  
 TCGA-06-1802\_family\_history.tsv

**Example of Open Science and reproducible research**  
**TCGA** **ICIA**

**is multidimensional approach?**

Data Access Detailed Description Clinicals & Data usage Policy Versions

Version 4 (Current): Updated 2020/05/29

Data Type Download all or Query/Filter

Images (DICOM, 73,058)

Tissue Slice Images (web)

Clinical Data (TXT)

Biomedical Data (TXT)

Genomics (web)

Detailed Description

Image Statistics

Mobilities MB

Number of Participants 262

Number of Studies 876

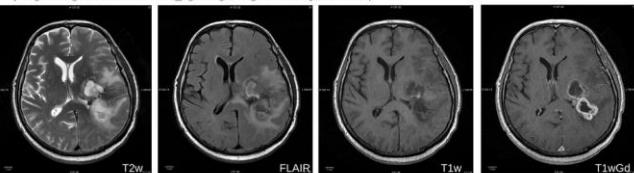
Number of Series 9,412

Number of Images 491,158

Images Size (GB) 73.5

Updated clinical data file with latest spreadsheets from GDC. Added new biomaterials

\* TCGA, TCGA-GBM, 09-18-2013.tsv → TCGA-06-1802.gdt, downloaded\_20200630\_211554\_634092.json.gz, TCGA-06-1802.zip



Arvid Ljønnevold, NordDoc 2022

# Plan



What *is* deep learning?



How do you *do* deep learning?



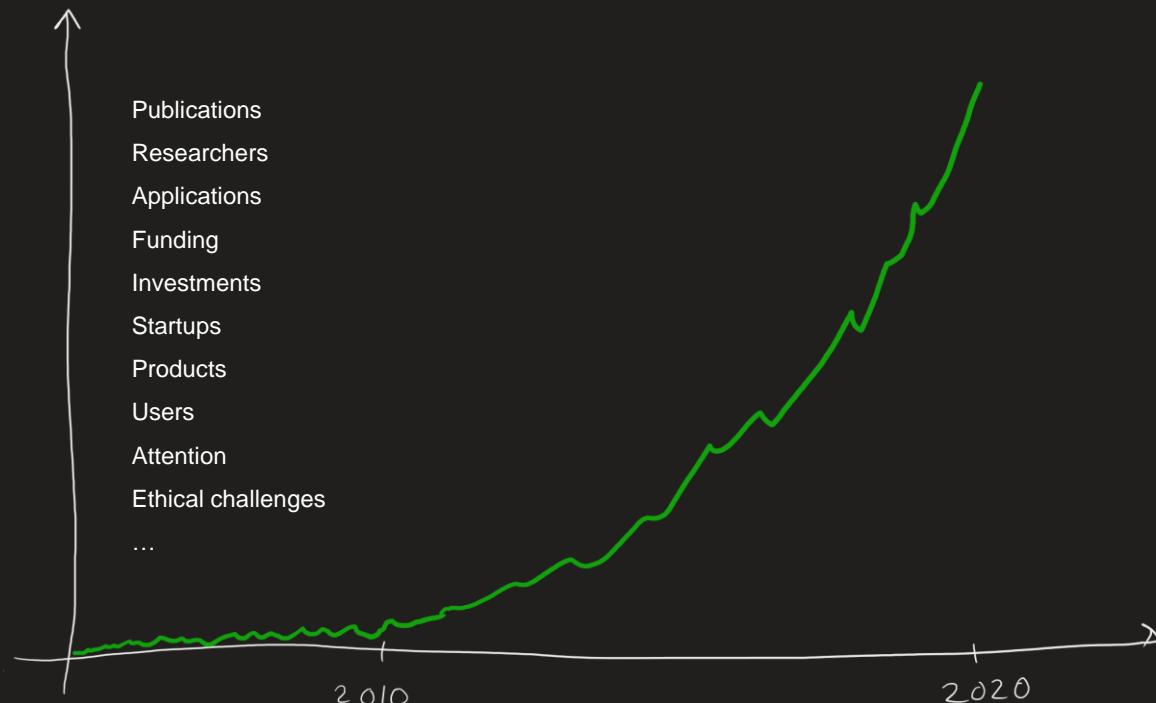
Some perspectives,  
opportunities and  
challenges

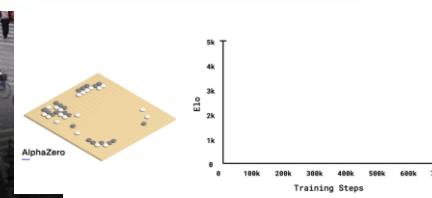
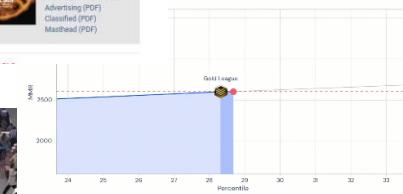
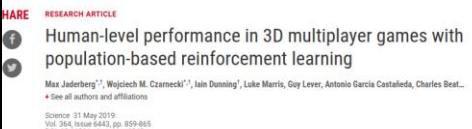
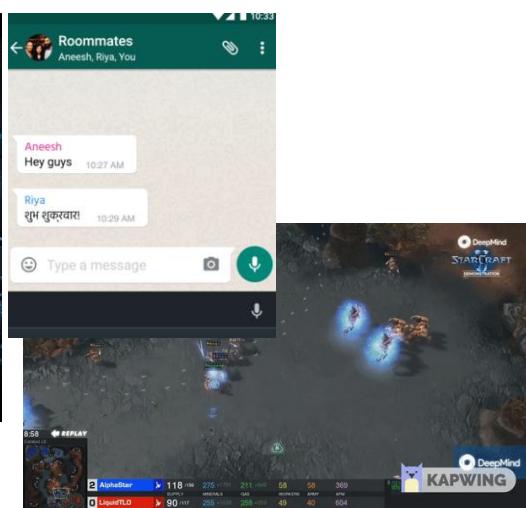
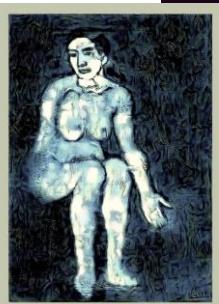
## Artificial intelligence

# Machine learning

# Deep learning

Representation learners







## Skin cancer can't hide from deep-learning diagnostics

May 21, 2019 | Dave Pearson | Diagnostics



NEWS CAREERS JOURNALS ▾

# Science



LOG IN

## FDA Approves AI Tool That Can Detect Wrist Fractures

The FDA has approved a new tool that can improve the accuracy of wrist fracture diagnosis.



By Jessica Miley

May 28th, 2018



## Deep learning improves detection of polyps during colonoscopy

November 07, 2018 | Matt O'Connor | Artificial Intelligence



Deep learning predicts OCT measures of diabetic macular thickening



By Steve Lenier

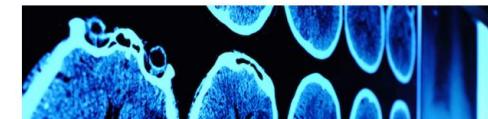
May 21, 2019

## Novel Molecules Designed by Artificial Intelligence May Accelerate Drug Discovery

**TOPICS:** Artificial Intelligence Biotechnology InSilico Medicine

## Deep Learning Tool Identifies Disease Markers of Alzheimer's

A deep learning tool was able to rapidly detect markers of Alzheimer's disease in human brain tissue.



AI cracks the code of protein complexes—providing a road map for disease

software maps thousands of the partnered proteins

NOV 2021 • 2:00 PM • BY ROBERT F. SERVICE

Results

### tein complex prediction with AlphaFold-Multimer

Richard Evans, Michael O'Neill, Alexander Pritzel, Natasha Antropova, Andrew Green, Augustin Zidek, Russ Bates, Sam Blackwell, Jason Yim, Olaf Sebastian Bodensteiner, Michał Zieliński, Alex Bridgland, Anna Potapenko, Andrew Athreya, Tatyana Savchenko, Rishabh Jain, Ellen Clancy, Pushmeet Kohli, John emis Hassabis  
<https://doi.org/10.1101/2021.10.04.463034>

This article is a preprint and has not been certified by peer review (what does this mean?).

NOVEMBER 15, 2017

## Stanford algorithm can diagnose pneumonia better than radiologists

Stanford researchers have developed a deep learning algorithm that evaluates chest X-rays for signs of disease. In just over a month of development, their algorithm outperformed expert radiologists at diagnosing pneumonia.



BY TAYLOR KUBOTA

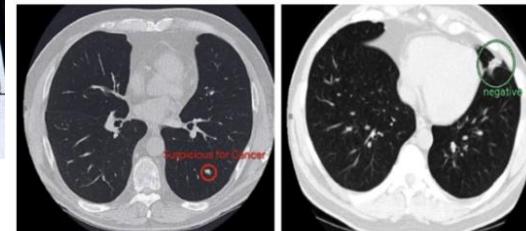
Stanford researchers have developed an algorithm that offers diagnoses based off chest X-ray images. It can diagnose up to 14 types of medical conditions and is able to diagnose pneumonia better than expert radiologists working alone. A paper about the algorithm, called CheXNet, was published Nov. 14 on the open-access, scientific preprint website arXiv.

"Interpreting X-ray images to diagnose pathologies like



## Google's lung cancer detection AI outperforms 6 human radiologists

KHARI JOHNSON @KHARIJOHNSON MAY 20, 2019 8:00 AM

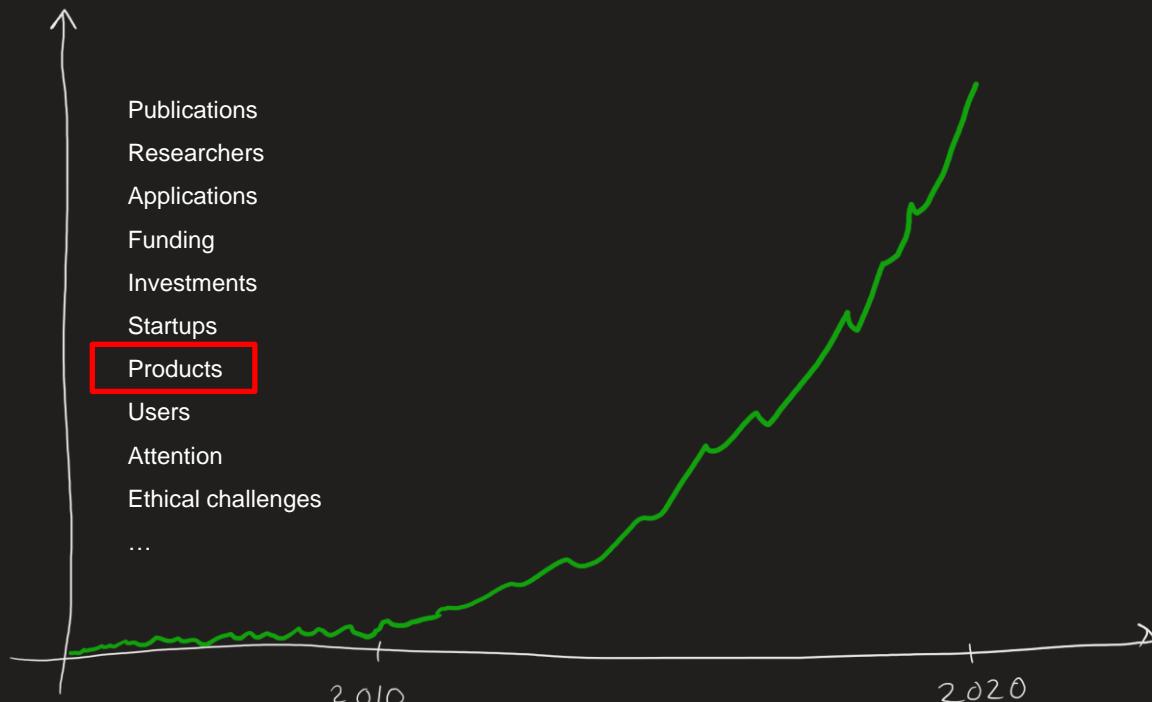


Artificial intelligence

Machine learning

**Deep learning**

*Representation learners*





All of these are available on the European market per May 2022 (CE marked)

SAMSUNG

Imaging Biometrics  
An IQ-AI Company



Cercare Medical

Visiana

ULTROMICS™

Motient

Perspectum

br

REVEA  
INFORMING PRECISION

CASIS  
Cancer Simulation & Imaging Software

MILVUE

ARTERYS

Riverain  
TECHNOLOGIES

aidoc

veolity

Quantib

mediaire

imagilys

thiron

Combinostics

GLEAMER

VIVOS®

Quibim  
Radiobotics  
Augmented Radiology

NICOLAB

Vara

cortech

deep01  
artificial intelligence

Behold.ai

BRAINLAB

Lucida Medical

SIEMENS

DIA

Perspectum

Imaging Biometrics  
An IQ-AI Company

The screenshot shows a search results page for "AI for Radiology". The results are filtered by "Modality: CT" and "CE class: I". There are 198 results. The products listed include:

- Avicenna AI CINA-ICH**: Subspecialty: Neuro, Modality: CT. CE: Class I - MDD, FDA: Class II.
- Avicenna AI CINA-LVO**: Subspecialty: Neuro, Modality: CT. CE: Class I - MDD, FDA: Class II.
- Avicenna AI CINA-ASPECTS**: Subspecialty: Neuro, Modality: CT. CE: Class I - MDD, FDA: Class II.
- Aidoc Brain aneurysm (BA)**: Subspecialty: Neuro, Modality: CT. CE: Class I - MDD, FDA: Class II.

## OBS: Medical Device Directive (MDD) versus Medical Device Regulation (MDR)

The background of the slide is a photograph of a mountainous region. In the foreground, there's a grassy slope with some rocks. A dirt path starts from the bottom left and winds its way up the hill. In the middle ground, there's a valley with more greenery and a few small buildings. The background features a range of mountains with sharp peaks, some of which have snow or ice on them. The sky is filled with large, white clouds.

Tasks

Modalities

Anatomical regions

...

# Narrow

Data availability

Competency availability

Medical feasibility

Technical feasibility

Organizational feasibility

Regulatory feasibility

Theoretical limitations

Incentives

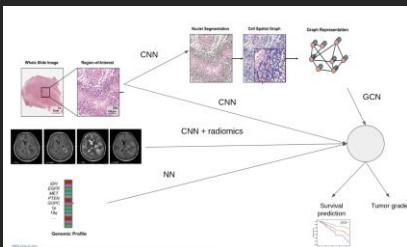
...

The research side is somewhat broader...

...trying to investigate how to integrate heterogeneous data about patients and disease processes.

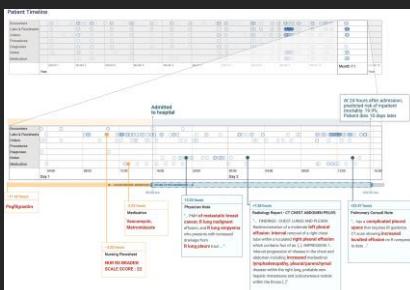
Two examples:

### Integrated diagnostics



Chen et.al. IEEE Trans Med Imaging, 2022, [10.1109/TMI.2020.3021387](https://doi.org/10.1109/TMI.2020.3021387)

### Deep learning and electronic health records



From EHR to length-of-stay and in-hospital mortality

Rajkomar et.al., NPJ Digit Med, 2018, [10.1038/s41746-018-0029-1](https://doi.org/10.1038/s41746-018-0029-1)

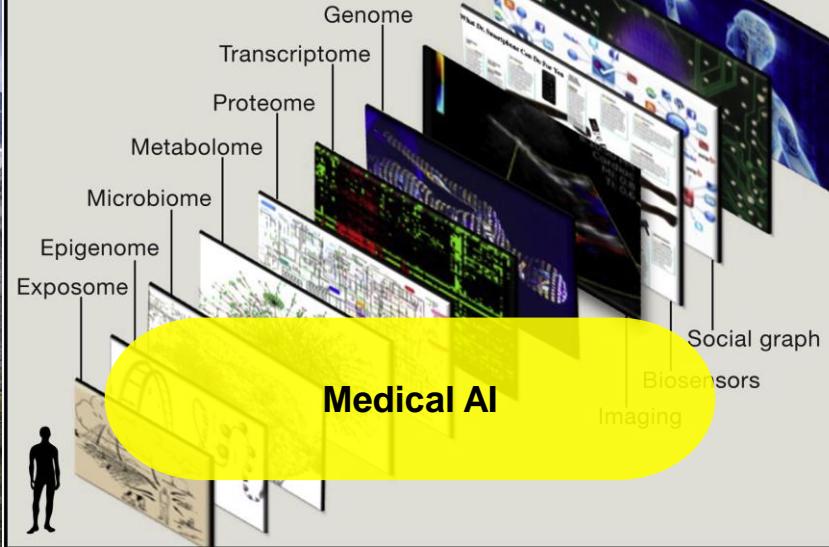


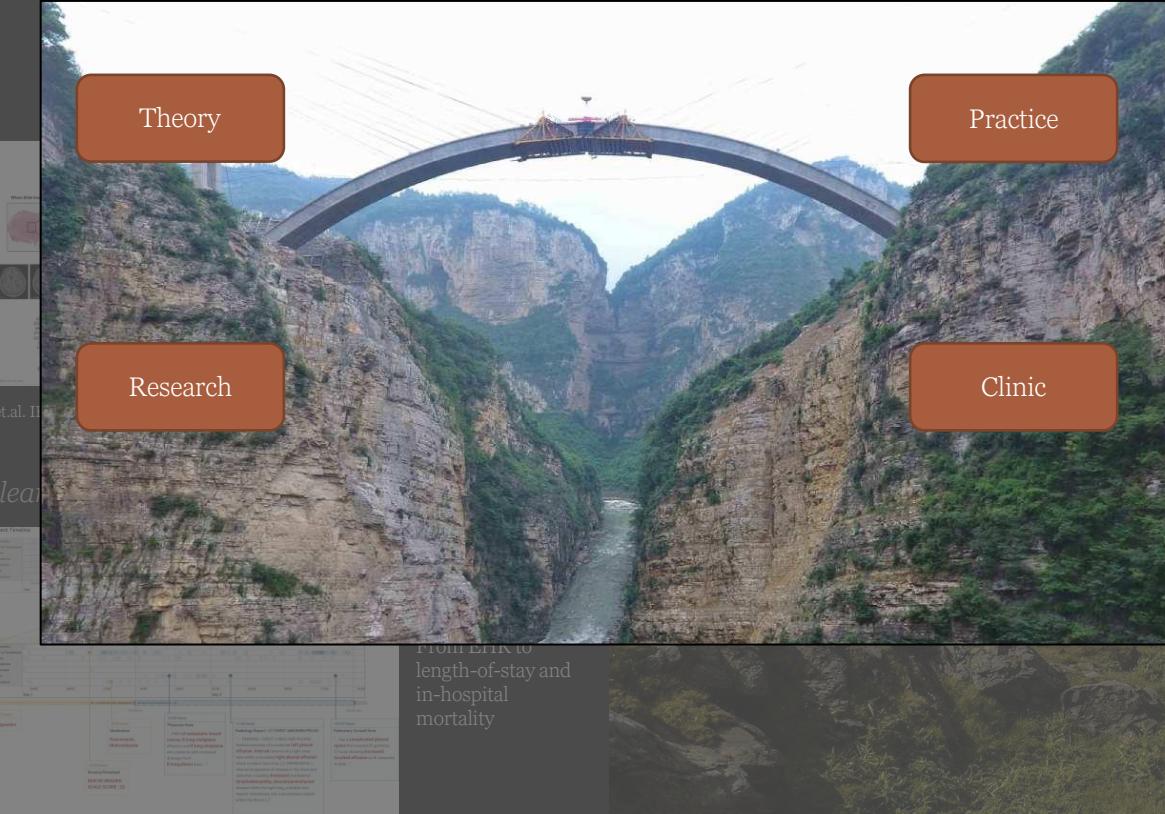
Illustration from E. Topol, Cell, 157(1), 2014



The research side is somewhat broader...

...trying to investigate how to integrate heterogeneous data  
about patients and disease processes.

Two examples:





Machine learning in medicine



A wide-angle photograph of a majestic mountain range under a bright blue sky with scattered white clouds. The mountains are rugged, with exposed rock faces and patches of green vegetation. In the foreground, there are several small, vibrant green lakes nestled in the valleys. A large, semi-transparent blue oval is positioned in the center-left of the image. Inside the oval, the text "Practical machine learning" and "&" is written in a white sans-serif font. Below that, the words "Machine learning engineering" are written in a smaller, italicized white sans-serif font.

Practical machine learning  
&  
*Machine learning engineering*



Machine  
learning



Machine learning

Machine learning



Machine  
learning



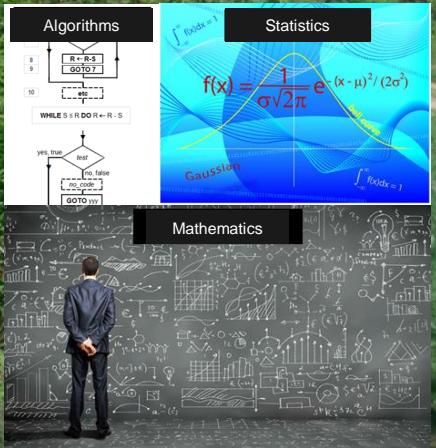
Machine learning  
models

Machine learning  
engineering

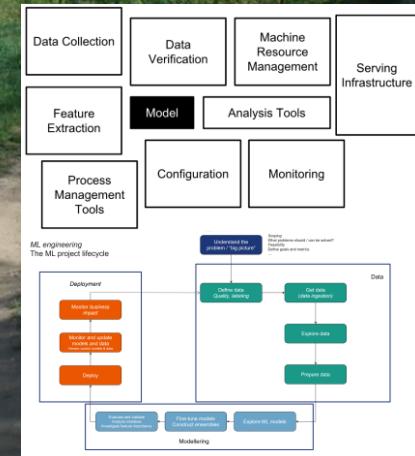


Machine learning

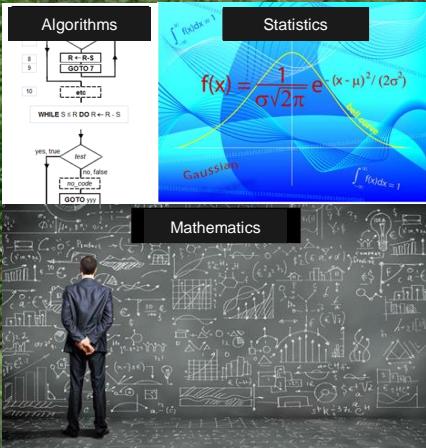
## Machine learning models



## Machine learning engineering



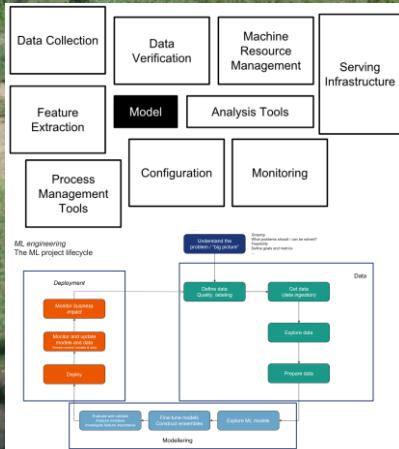
## Machine learning models



Machine  
learning

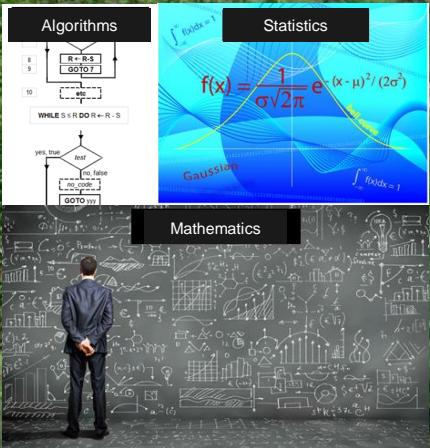


## Machine learning engineering

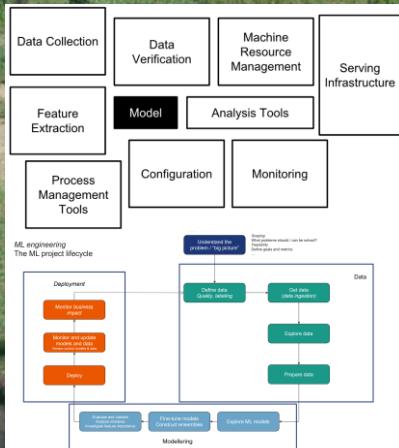




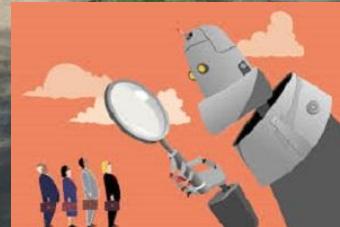
# Machine learning models



# Machine learning engineering



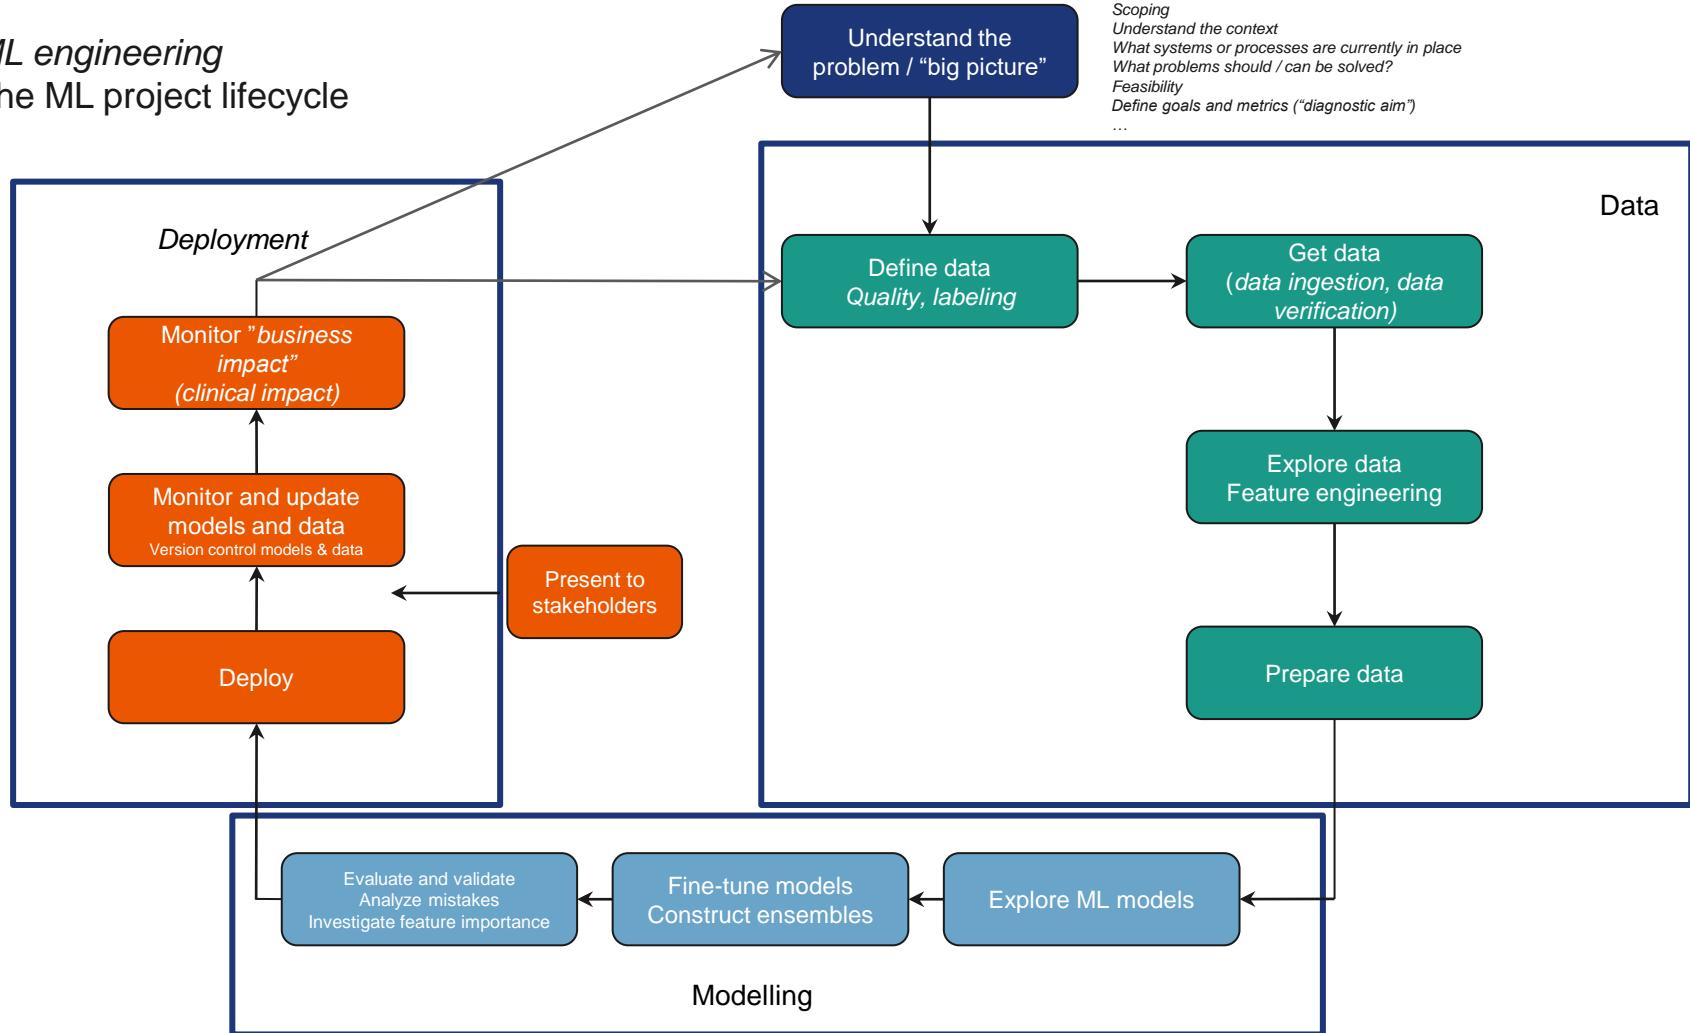
# Data and society



How do you make a machine learning-based solution?

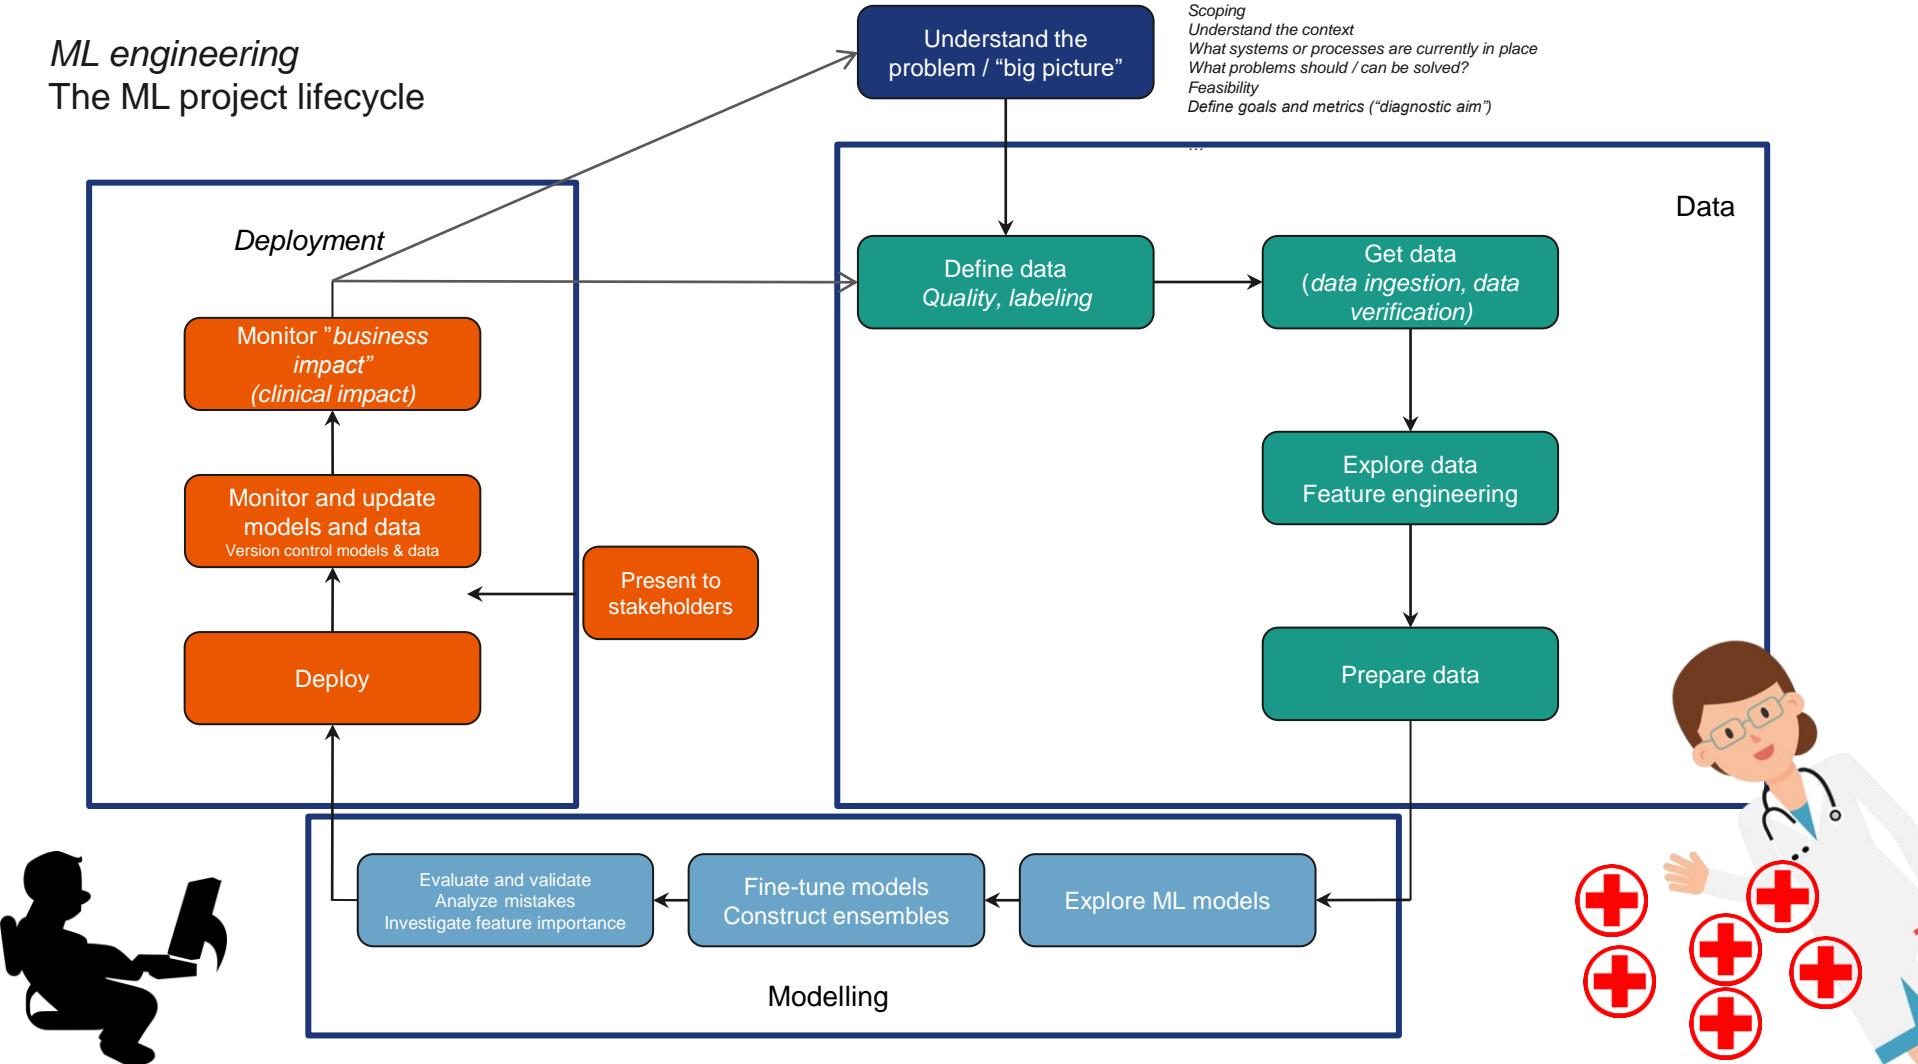
# ML engineering

## The ML project lifecycle

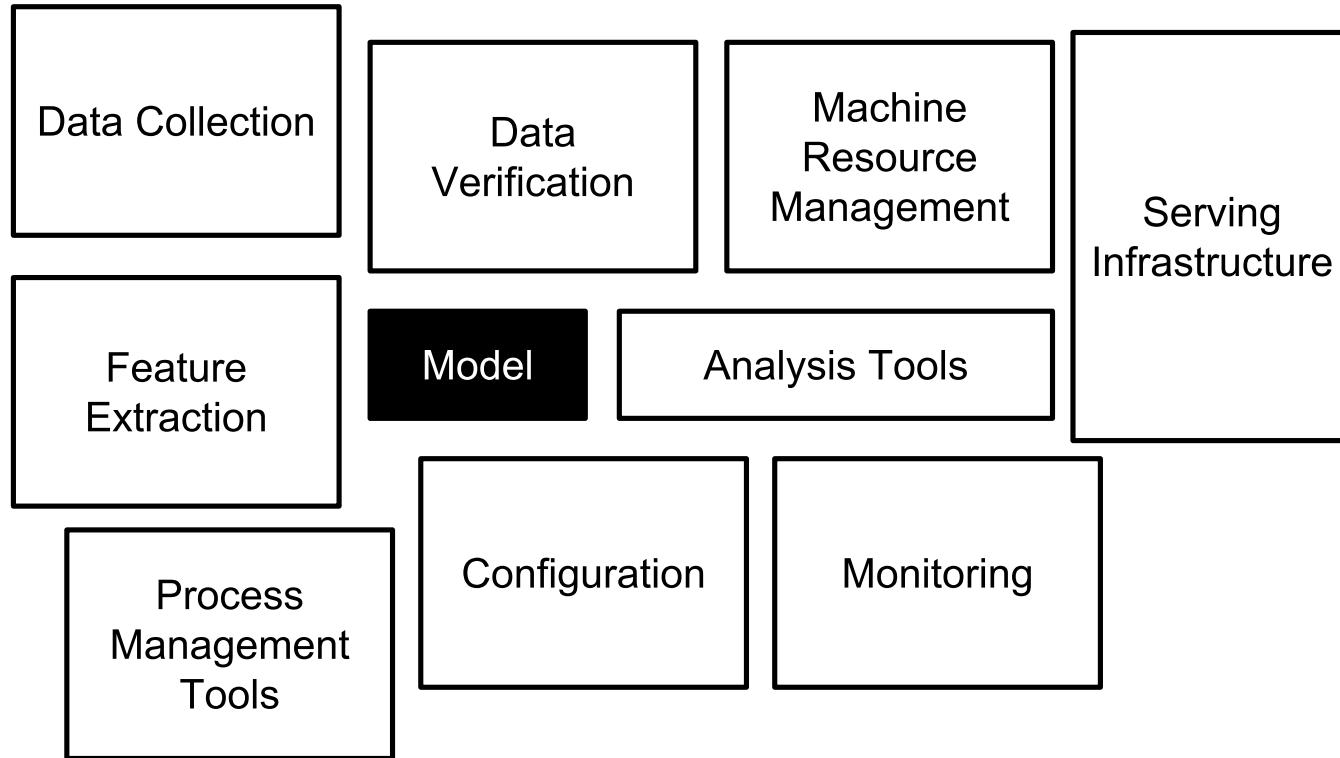


# ML engineering

## The ML project lifecycle



Model



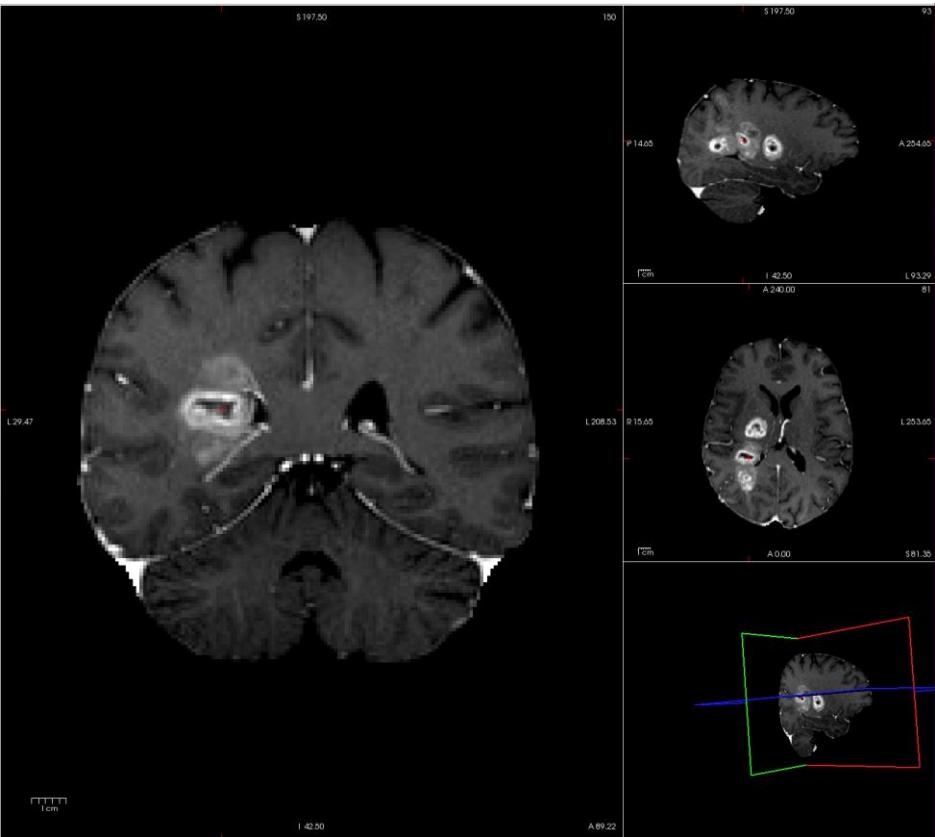


The background image shows a steep, green-covered mountain slope. A light-colored, winding path or stream bed cuts through the vegetation. The terrain is rocky and uneven, with patches of grass and small trees. The lighting suggests a bright, sunny day.

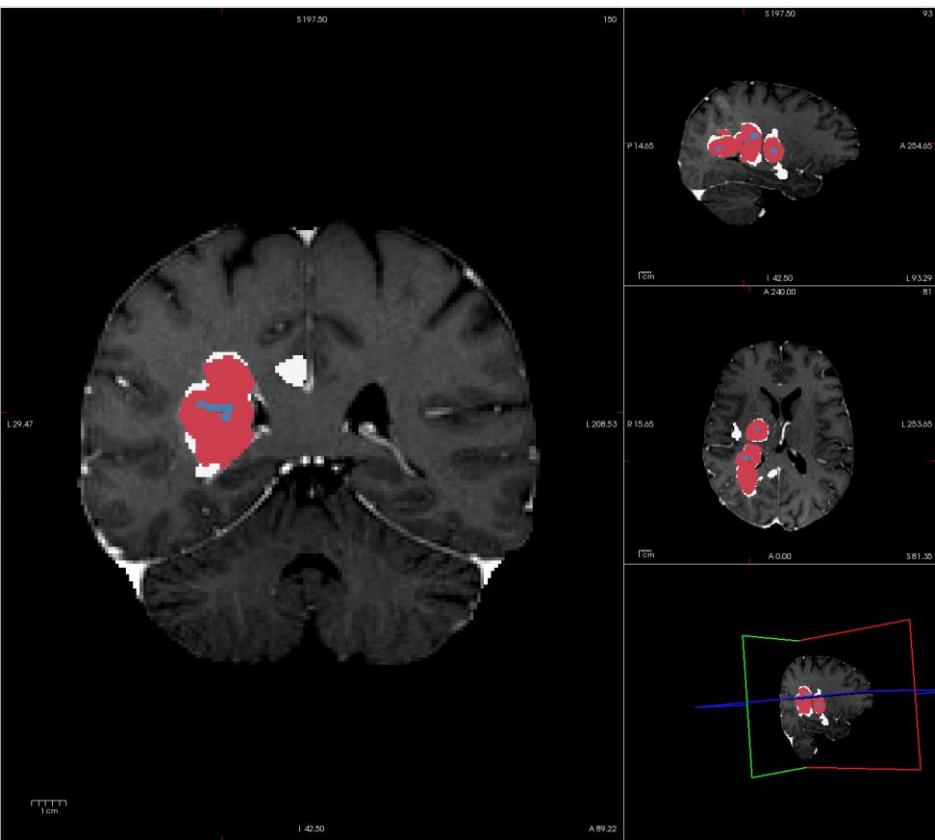
Machine learning in medicine

Model

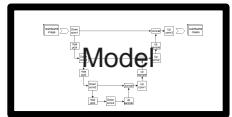
Model



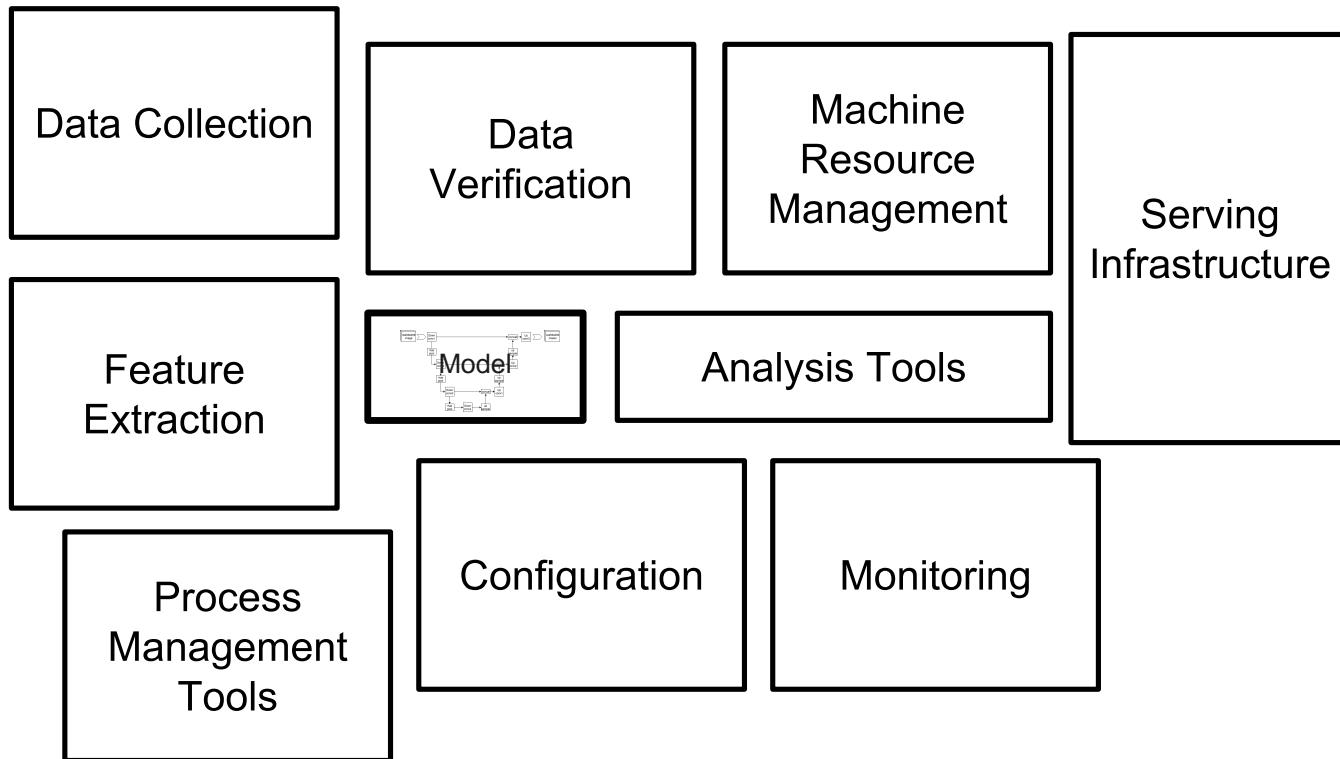
Model







In radiology, solutions must be compatible with existing infrastructure!





Hauke Bartsch  
Computer Science



Zhanbolat  
Satybaldinov  
Software  
engineering

In radiology, solutions must be compatible with existing infrastructure.

# Research Information System for the Western Norway Regional Health Authorities

medical research data in Helse Førde, Bergen, Fonna, and Stavanger

## Partners

Researchers – Mohn Medical Imaging and Visualization  
Healthcare professionals – Radiology  
Technologists – Helse Vest IKT  
Funding – The Research Council of Norway

## Features

Data migration, anonymization, exchange, and data processing  
Commercial image archiving and viewing platform  
Electronic data capture (eCRF)

## Hospital infrastructure for research

Connects 4 major hospitals and 30 treatment centers



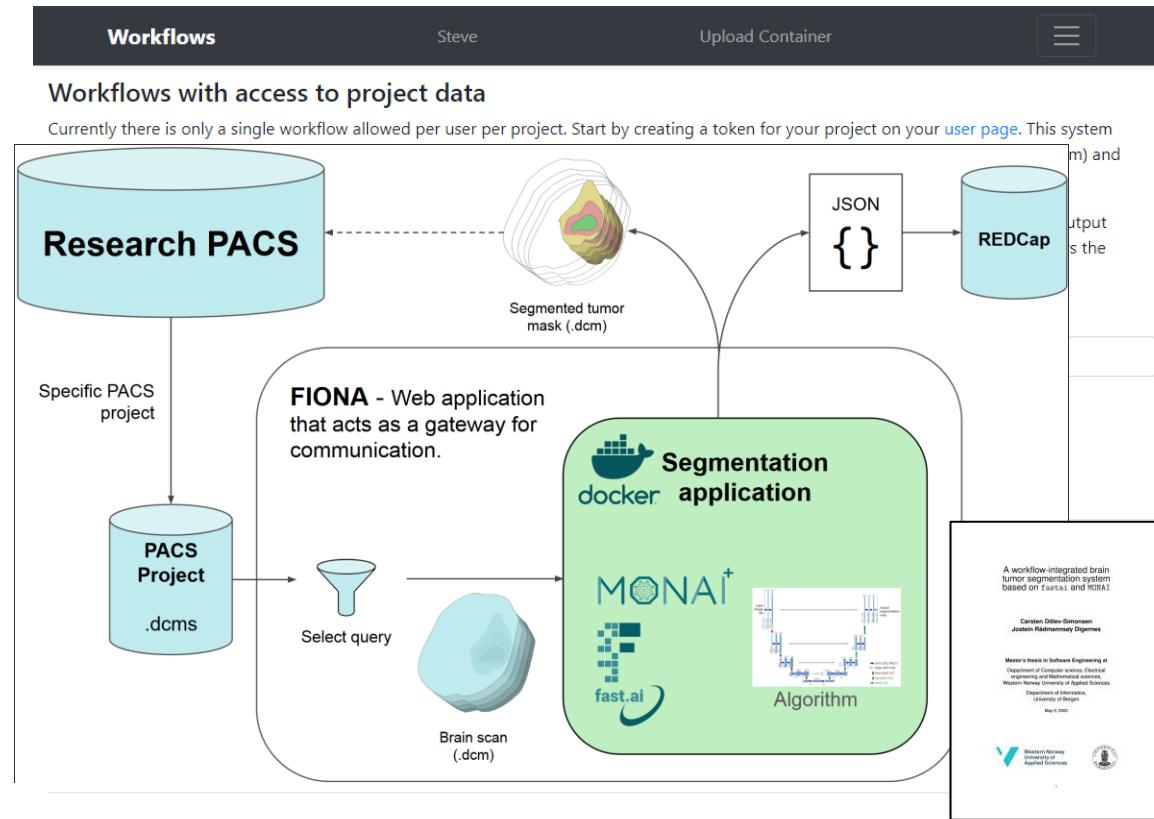
<https://github.com/mmiv-center/Research-Information-System>

# AI workflows

User driven activity supporting deep learning on medical images with training and prediction inside the hospital system.

Open-source code development see:

[github.com / mmiv-center / Research-Information-System / tree / master / components / Workflow-Image-AI](https://github.com/mmiv-center/Research-Information-System/tree/master/components/Workflow-Image-AI)



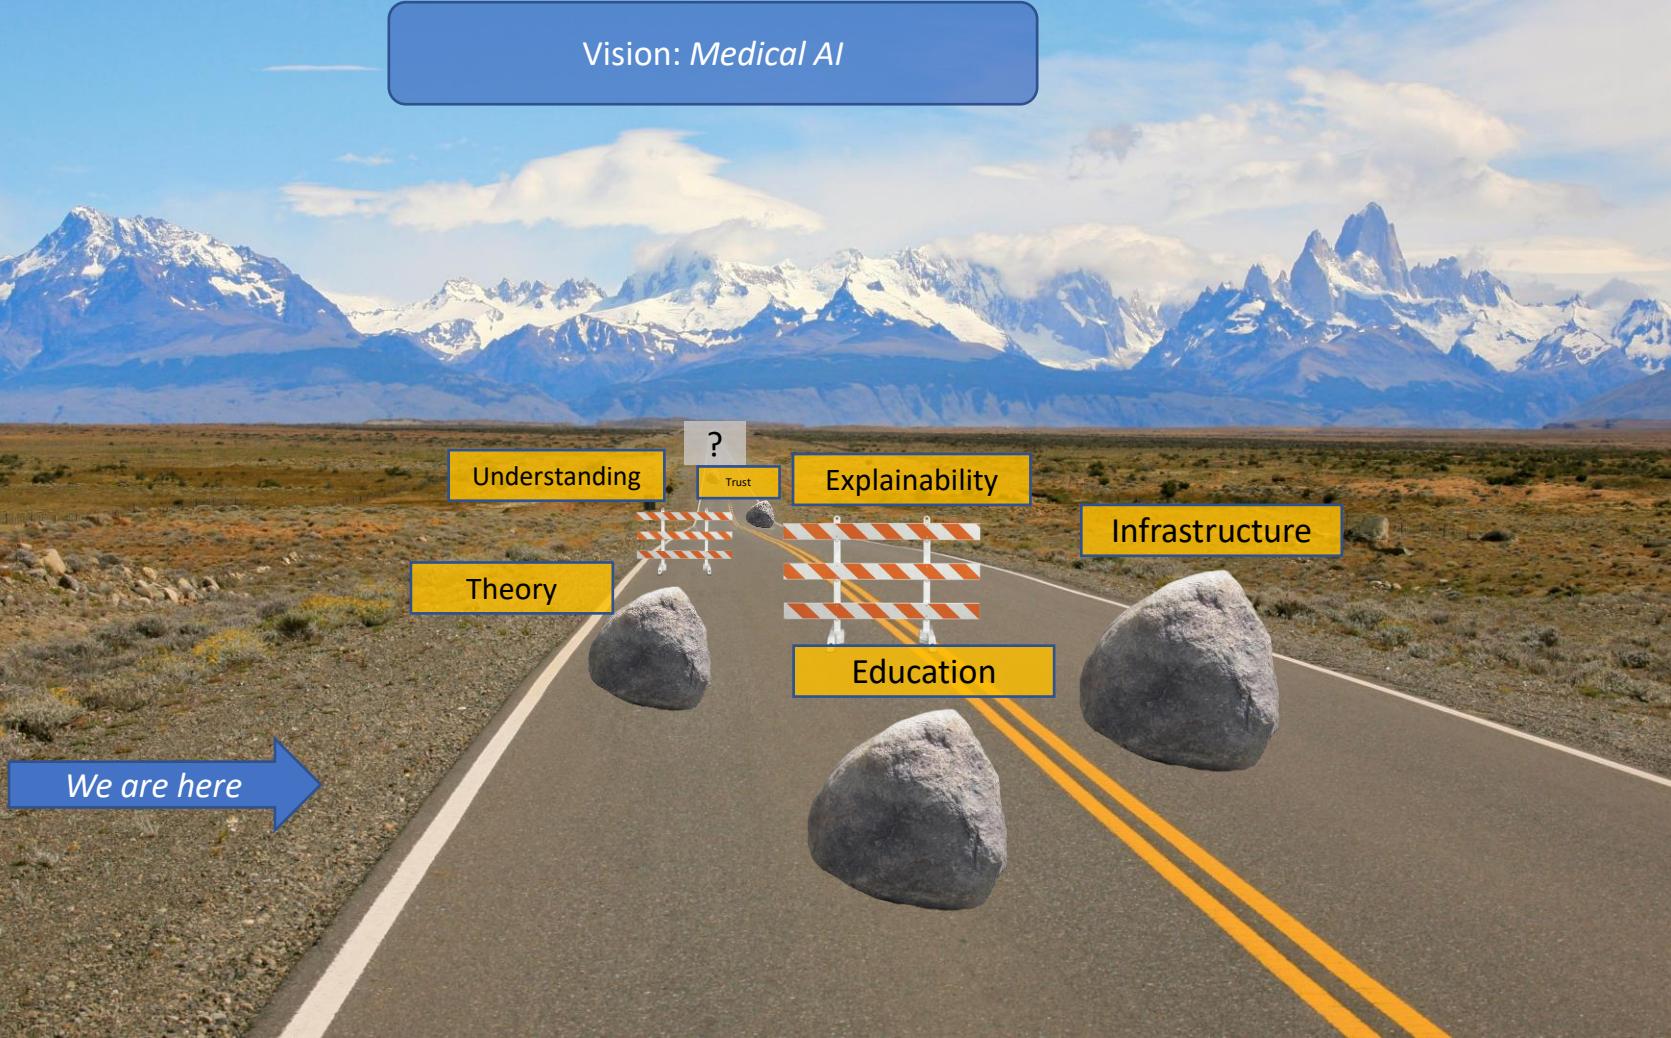
### Job list

The job list populates after a successful select statement has been provided. Based on the level (project|participant|study|series) a processing job might receive more or less data. The buttons allow a job to be scheduled. Processing is done in order of scheduling based on the compute resources allocated to the user.

At the beginning of the next section distributions for all detected output measures are displayed using luminance and opacity to indicate bins with many values. In order to be able to rank the datasets based on the *normality* of the calculated measures a numeric value and a color code is assigned to each job. Under certain restriction this rank can be used to identify possibly problematic cases for review.



## Vision: *Medical AI*



# The development of Medical AI should happen with medical experts in the driver's seat \*

in tight collaboration with (all the right kind of) scientists  
and technologists

We are here →

# The development of Medical AI should happen with medical experts in the driver's seat \*

in tight collaboration with (all the right kind of) scientists  
and technologists

A blue arrow pointing to the right, containing the text "We are here".

We are here



Vision: *Medical AI*

# The development of Medical AI should happen with medical experts in the driver's seat \*

in tight collaboration with (all the right kind of) scientists  
and technologists



We are here

infrastructure

Education